

Chapter 5. ROP

(Return-oriented Programming)

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering

Sogang University

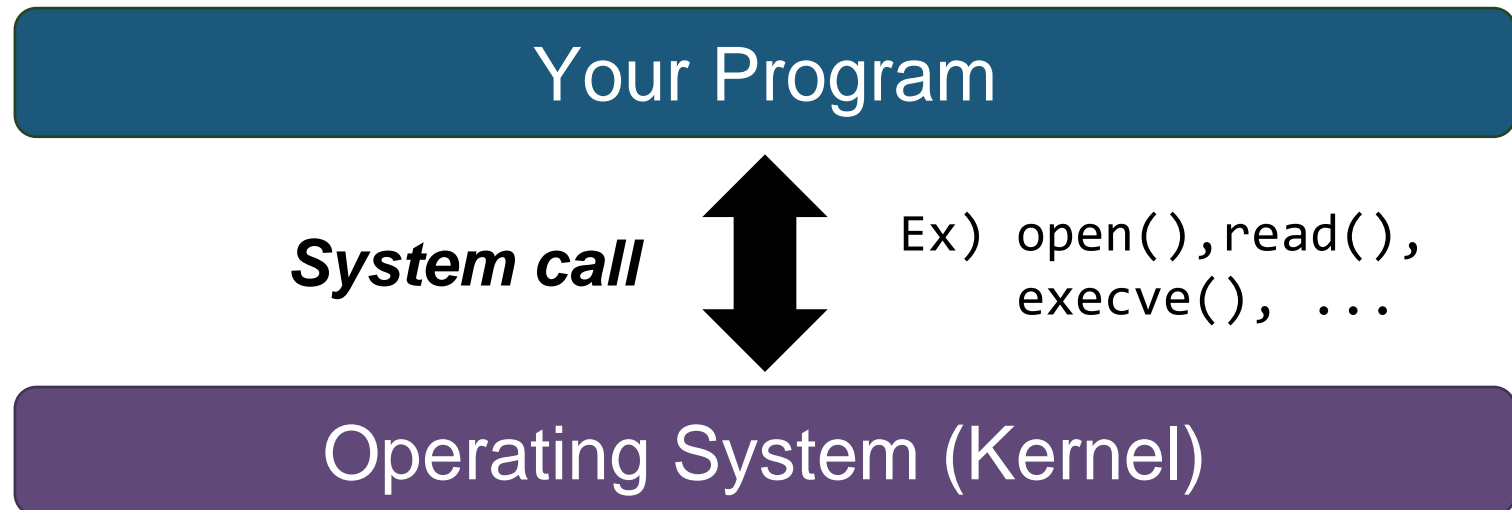
Topics

- **Background: system call**
- **The second round of war between attacker vs. defender**
 - Code reuse attack (**ROP**) to bypass mitigation with **ASLR**

System Call (Syscall)

■ Assume that you have written a C program that opens and reads a file

- What kind of x86-64 assembly instruction can we use?
- No instruction is solely reserved for opening or reading a file
- Instead, you must make some **request to OS (system call)**
 - The OS will do the task for you and return the result



Assembly Code For System Call

- It's similar to function call, but the function is in kernel
- When you setup particular registers properly and execute **syscall** instruction, system call is invoked
- What does actually happen during the system call?
 - Take *System Programming* or *Operating System* course

```
...  
mov    $0, %rsi      # %rsi must contain flag (option)  
mov    ..., %rdi     # %rdi must point to filename string  
mov    $0x2, %rax    # System call ID of open() is 2  
syscall
```

System Call Wrapper

- You can also invoke system calls like `open()`, `read()` or `execve()` in your C source code
 - You are actually calling a **wrapper** function around the `syscall`

Your Program

```
...  
inf fd = open("a.txt")  
...
```

```
...  
call 0x2000 # <open()>  
...
```

Library

```
<open syscall wrapper>  
0x2000: ...  
...  
0x200a: mov $0x2, %rax  
0x200f: syscall  
...
```

High-Level Library Functions

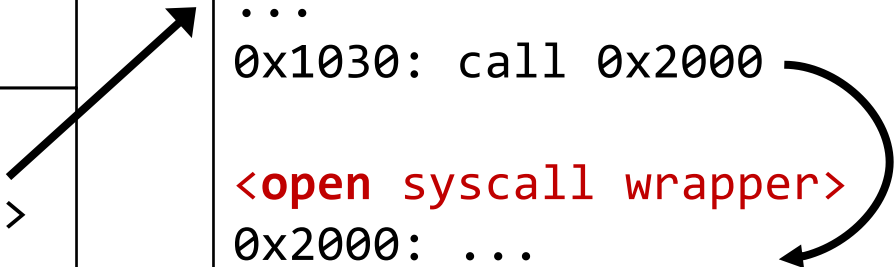
- You are probably more familiar with higher-level functions like `fopen()`, `fgets()`, `fread()`, etc.
 - Such functions are implemented by using system calls internally

Your Program

```
...  
File *f = fopen("a.txt")  
...  
  
...  
call 0x1000 # <fopen()>  
...
```

Library

```
<fopen() function>  
0x1000: ...  
...  
0x1030: call 0x2000  
  
<open syscall wrapper>  
0x2000: ...
```



The diagram illustrates the internal implementation of a high-level library function. An arrow points from the `call 0x1000 # <fopen()>` instruction in the program to the `0x1000: ...` entry in the library's `<fopen() function>`. Inside the library function, another arrow points from `0x1030: call 0x2000` to the `0x2000: ...` entry in the `<open syscall wrapper>`, showing how the high-level function is implemented using a lower-level syscall wrapper.

Topics


- Background: system call
- **The second round of war between attacker vs. defender**
 - Code reuse attack (**ROP**) to bypass mitigation with **ASLR**

Code Reuse Attack

- **Review:** By introducing NX (non-executable memory), **attack with shellcode** was effectively mitigated
 - Injected shellcode cannot be executed anymore
- **However, the attackers found a different way:**
 - **Code reuse attack:** use the existing code to achieve the goal
 - Ex) By executing `execve()` function in C library (`libc`) with `"/bin/sh"` as argument, the hacker can spawn a shell
 - Library code **must be executable**, so NX cannot prevent this

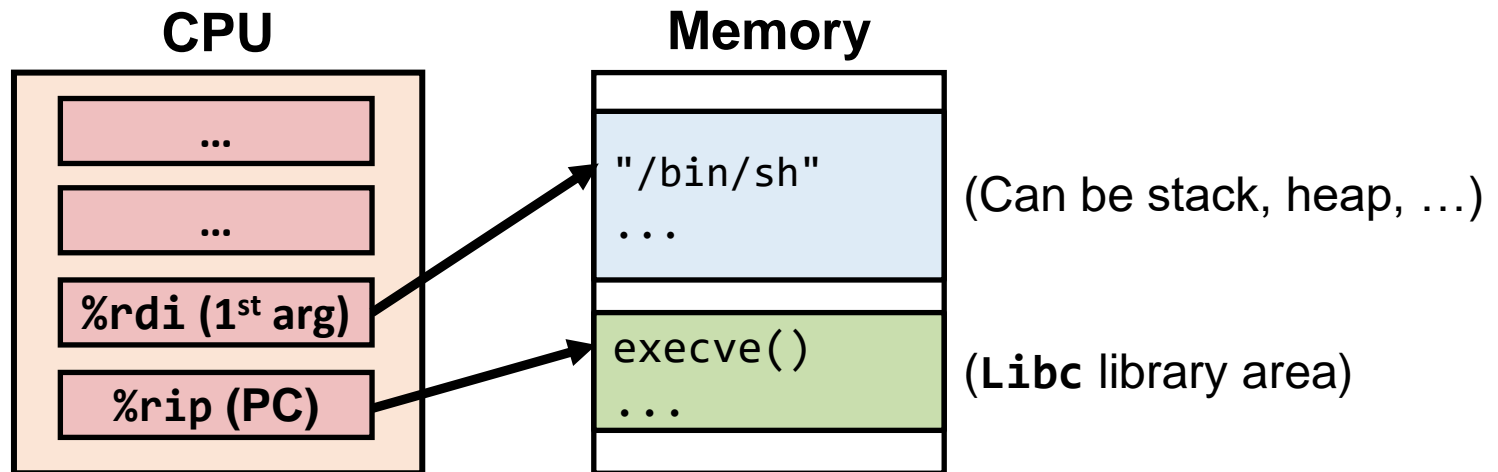
Return-to-Libc Attack

- Before 2000s, code reuse attack was easy
 - 32bit x86 system was the mainstream: according to its calling convention, **controlling the argument was easy**
 - Also, library code was placed in easily **predictable address**
- This primitive form of attack was called *return-to-libc*
 - But we will not cover this attack deeply (**outdated**)
 - Instead, we will focus on **x86-64 systems!**




Challenge #1

- In the calling convention of x86-64, **manipulating the arguments** of a function is difficult
- Even if we hijack the control to `execve()`, how can we pass `"/bin/sh"` argument to this function?
 - We must manipulate `%rdi` register to point to string `"/bin/sh"`
 - We could corrupt `%rip` register, but what about other registers?



Challenge #2

- In the early 2000s, **address space layout randomization (ASLR)** mitigation was introduced
 - Each memory section address is randomized **per execution**
 - However, for performance reason, the **Code** and **Data** sections were not randomized until recently




**We can overcome the first challenge
with a code-reuse attack called **ROP**
(return-oriented programming)**

ROP Gadget

- ROP attack uses small code chunks called **gadget**
 - Instruction sequence that **ends with ret** instruction
 - Your program often includes many unintended gadgets
- Consider **mov \$0xc35f, %rax** instruction below
 - What if we jump to **0x1003**, which is middle of this instruction?
 - Bytes are interpreted as instruction sequence **pop %rdi; ret**

Addr	Bytes	Instruction
0x1000	48 c7 c0 5f c3 00 00	mov \$0xc35f, %rax
0x1007	...	



Re-interpret

Addr	Bytes	Instruction
0x1003	5f	pop %rdi
0x1004	c3	ret


What can we do with gadgets?

■ Assume function `vuln()` with a buffer overflow

- BOF allows us to overwrite the saved return address
- Assume no unused space between `buf[]` and return address

```
void vuln(void) {  
    char buf[8];  
    gets(buf); // Buffer overflow  
}
```

(Before the BOF. Assume each block is 8-byte)




What can we do with gadgets?

- We will overwrite the saved return address like below
 - Note that we have corrupted **beyond** the saved return address

```
<vuln>
...
0x400875      ret      # Last instruction

... (some other function)
0x4009b3      pop %rdi  # Gadget to use
0x4009b4      ret
```

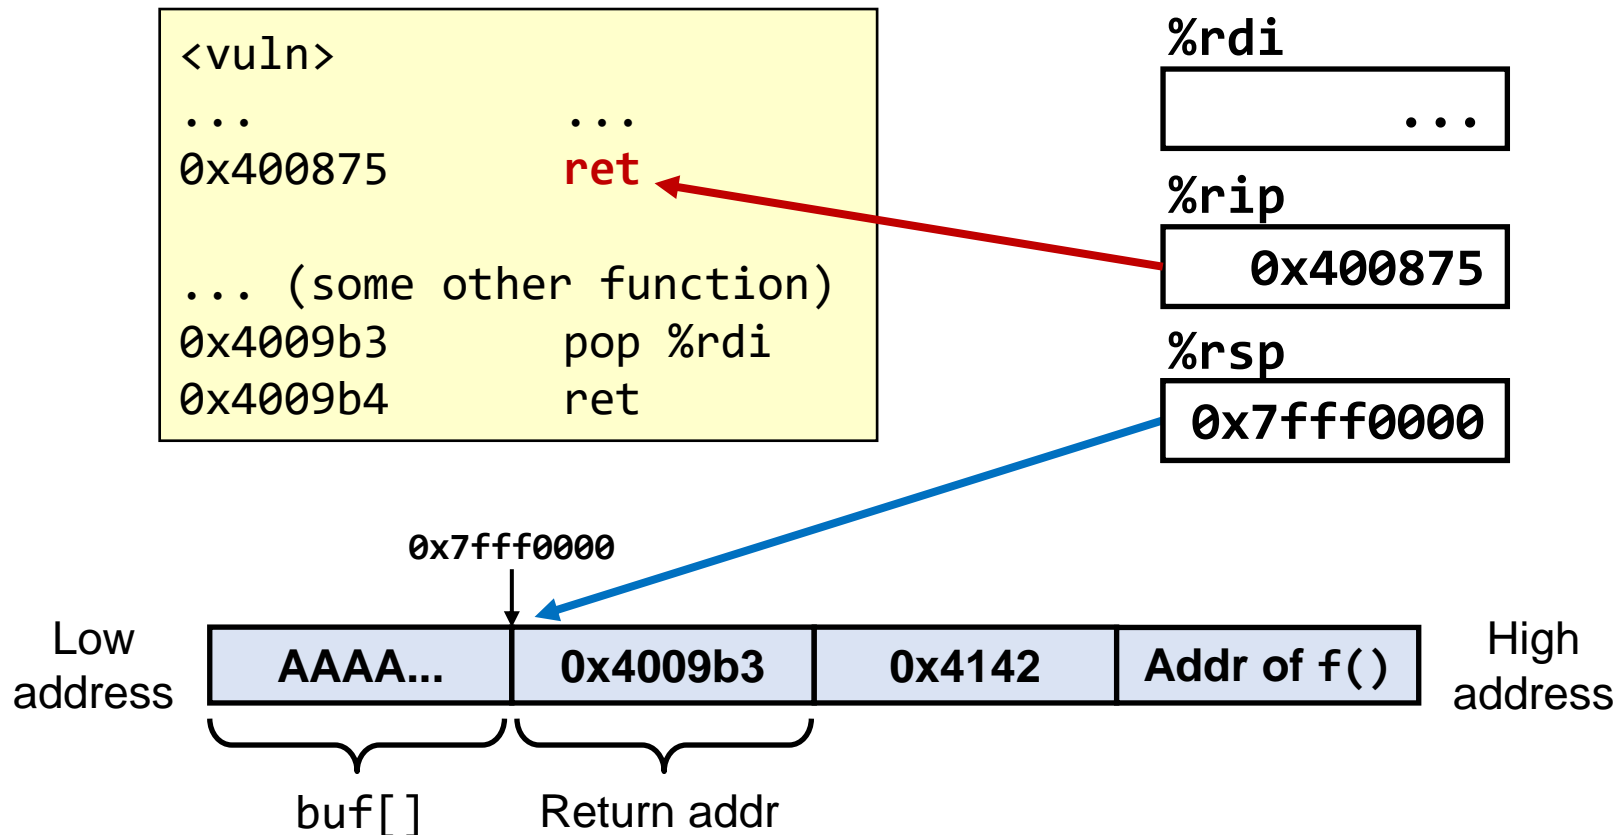
(Corrupted with BOF. $f()$ is arbitrary function we choose)



What can we do with gadgets?

■ Now, assume `vuln()` is about to return with **ret**

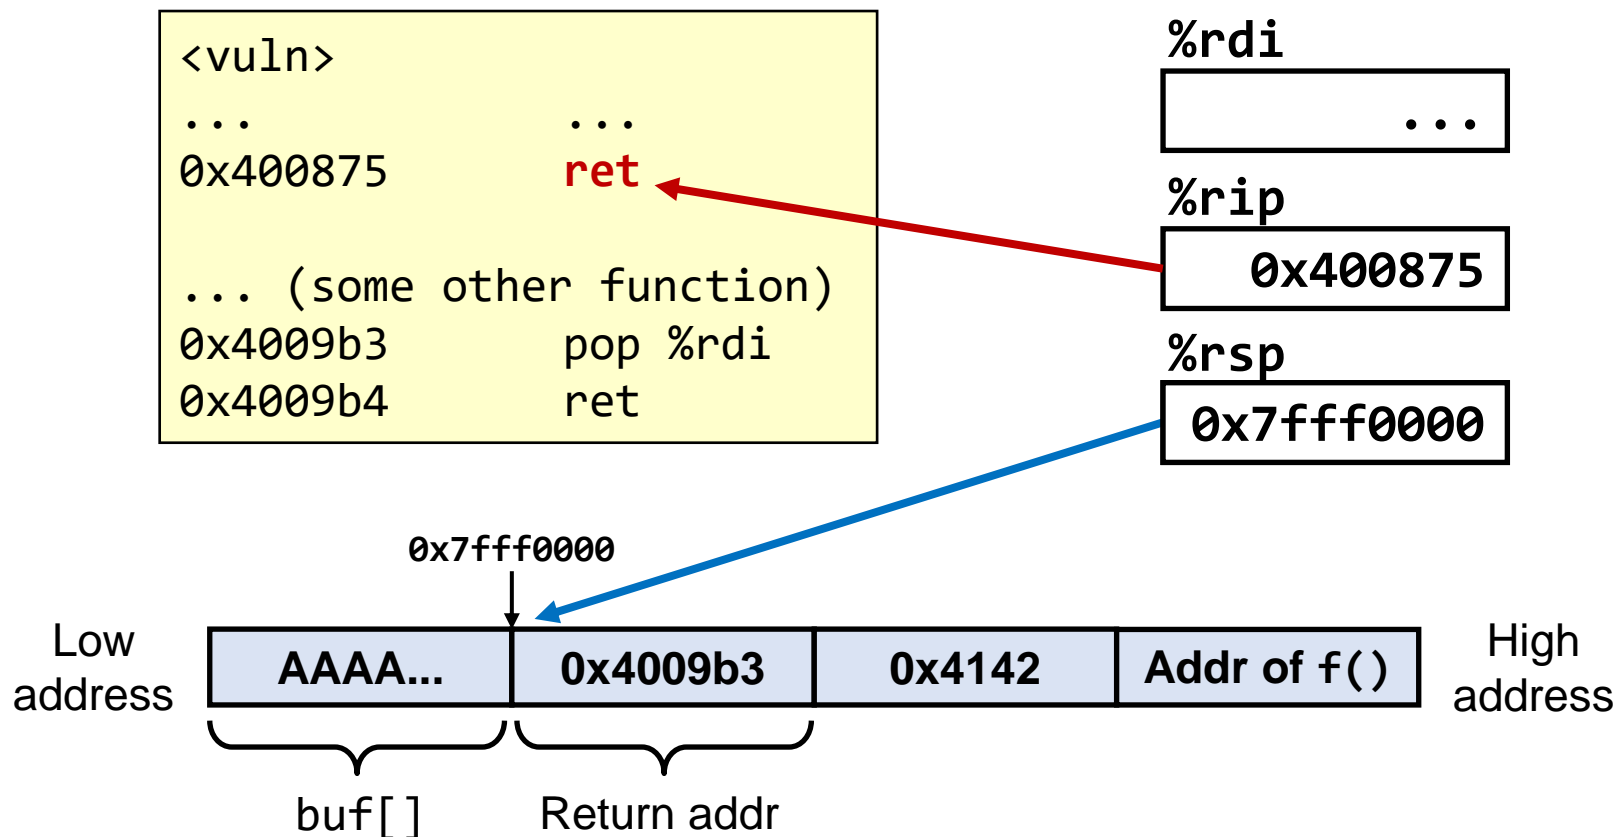
- Let's see what happens at this moment, step by step



What can we do with gadgets?

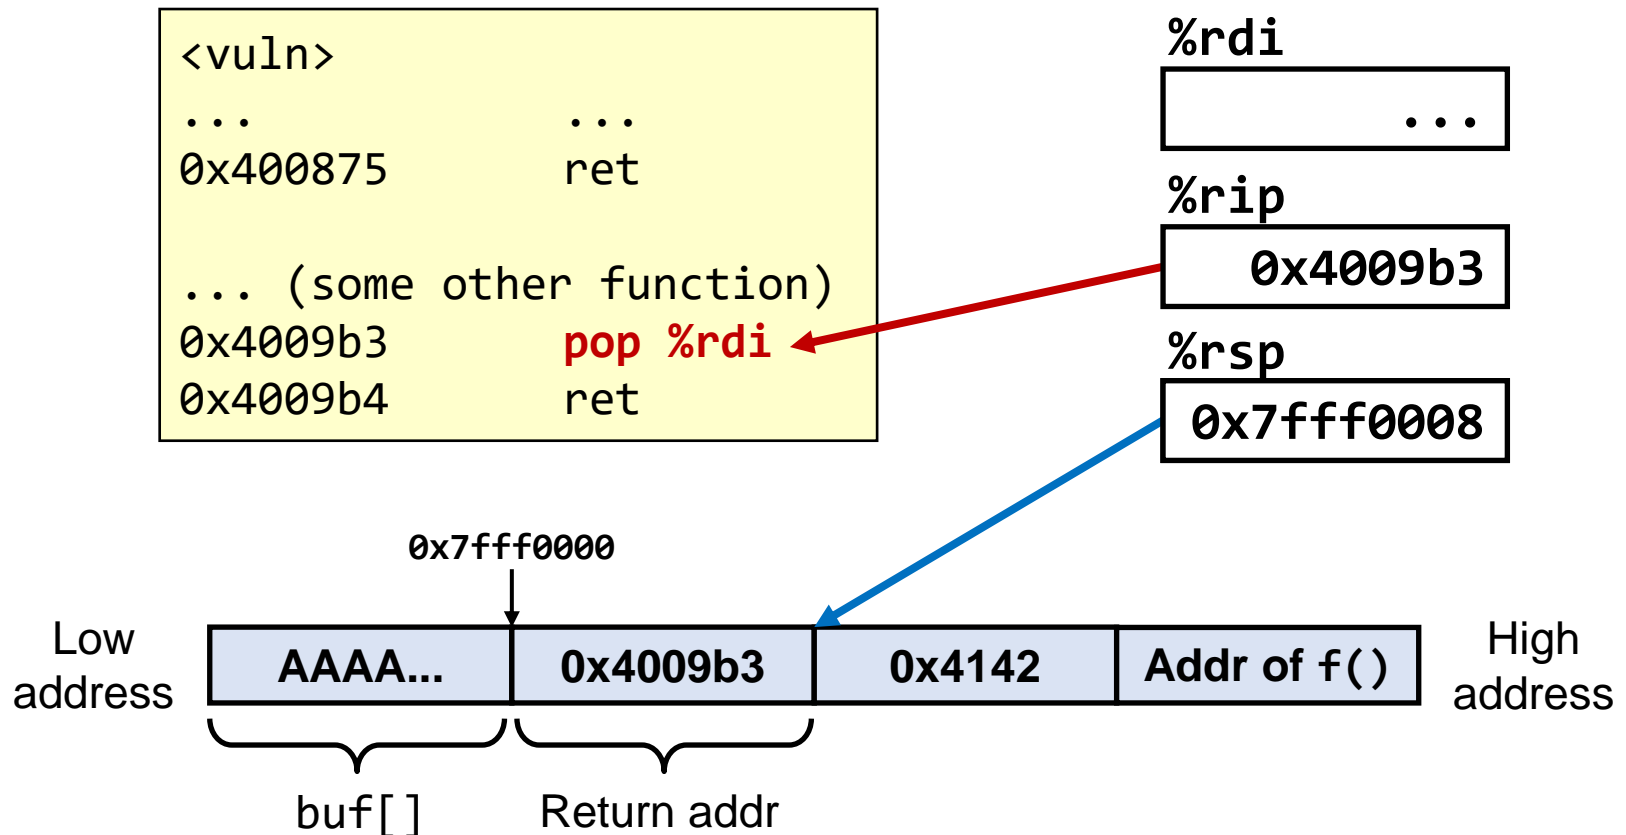
■ First, **ret** of `vuln()` will pop `0x4009b3` into `%rip`

- At the same time, `%rsp` will be incremented by 8



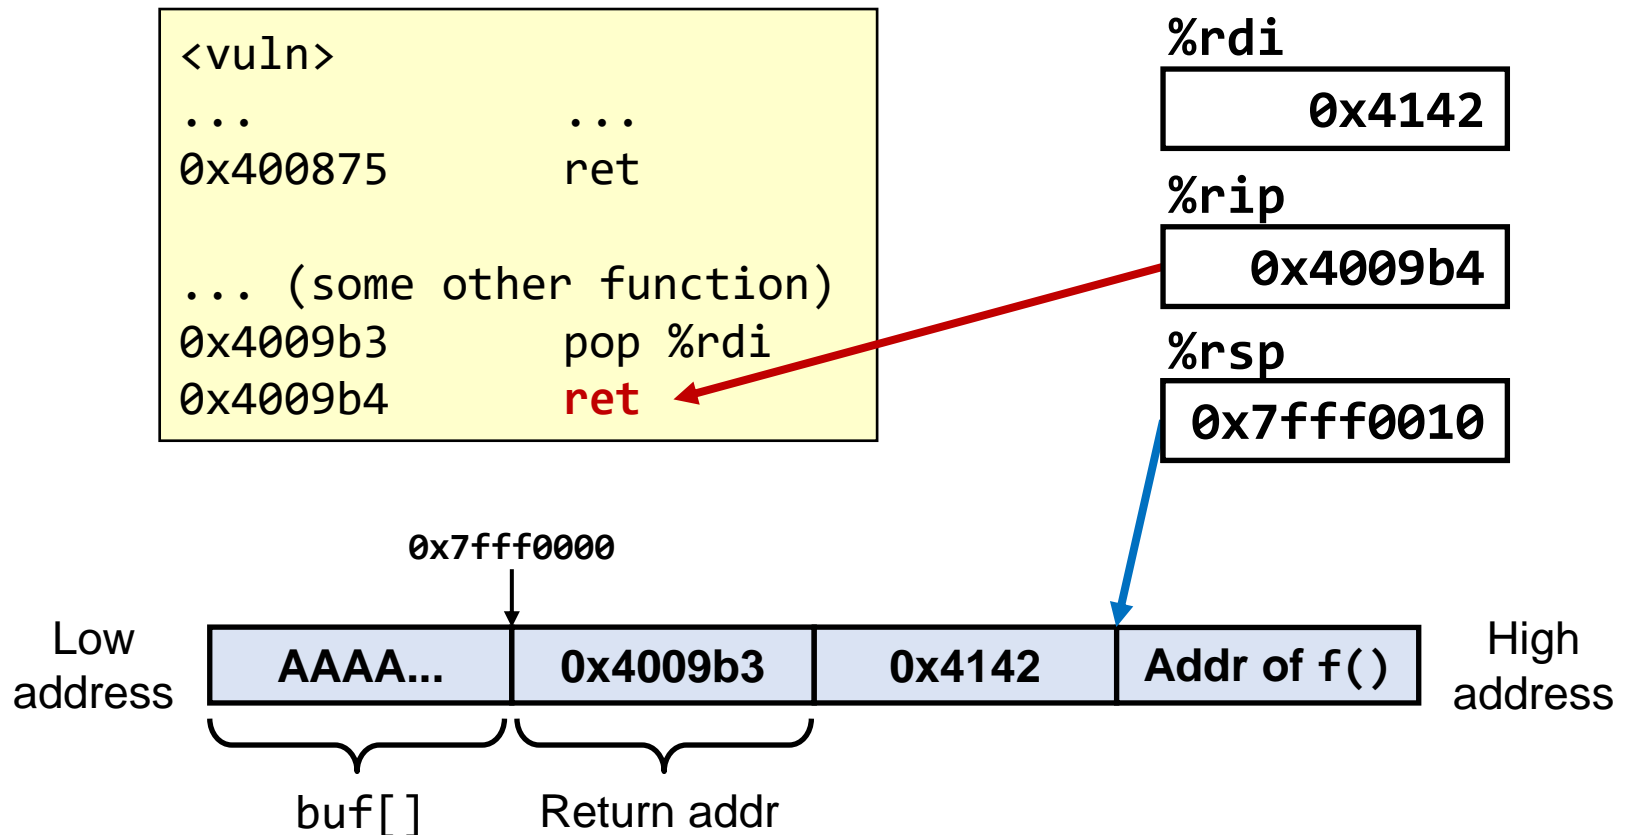
What can we do with gadgets?

- Then, `%rip` points at **pop %rdi** instruction in gadget
 - This instruction will pop `0x4142` into `%rdi` (= 1st argument)



What can we do with gadgets?

- Now, `%rip` is pointing at **ret** instruction in gadget
 - This will change `%rip` to point at the address of `f()`



We have just executed `f(0x4142)`!


ROP gadgets let us execute a chosen function `with arbitrary argument`

ROP Chain: Multiple Arguments

- By chaining multiple ROP gadgets, it is also possible to call a function with multiple arguments
- Let's overwrite saved return address and use 2 gadgets
 - **Caution:** Each block (even the small one) still represents 8-byte

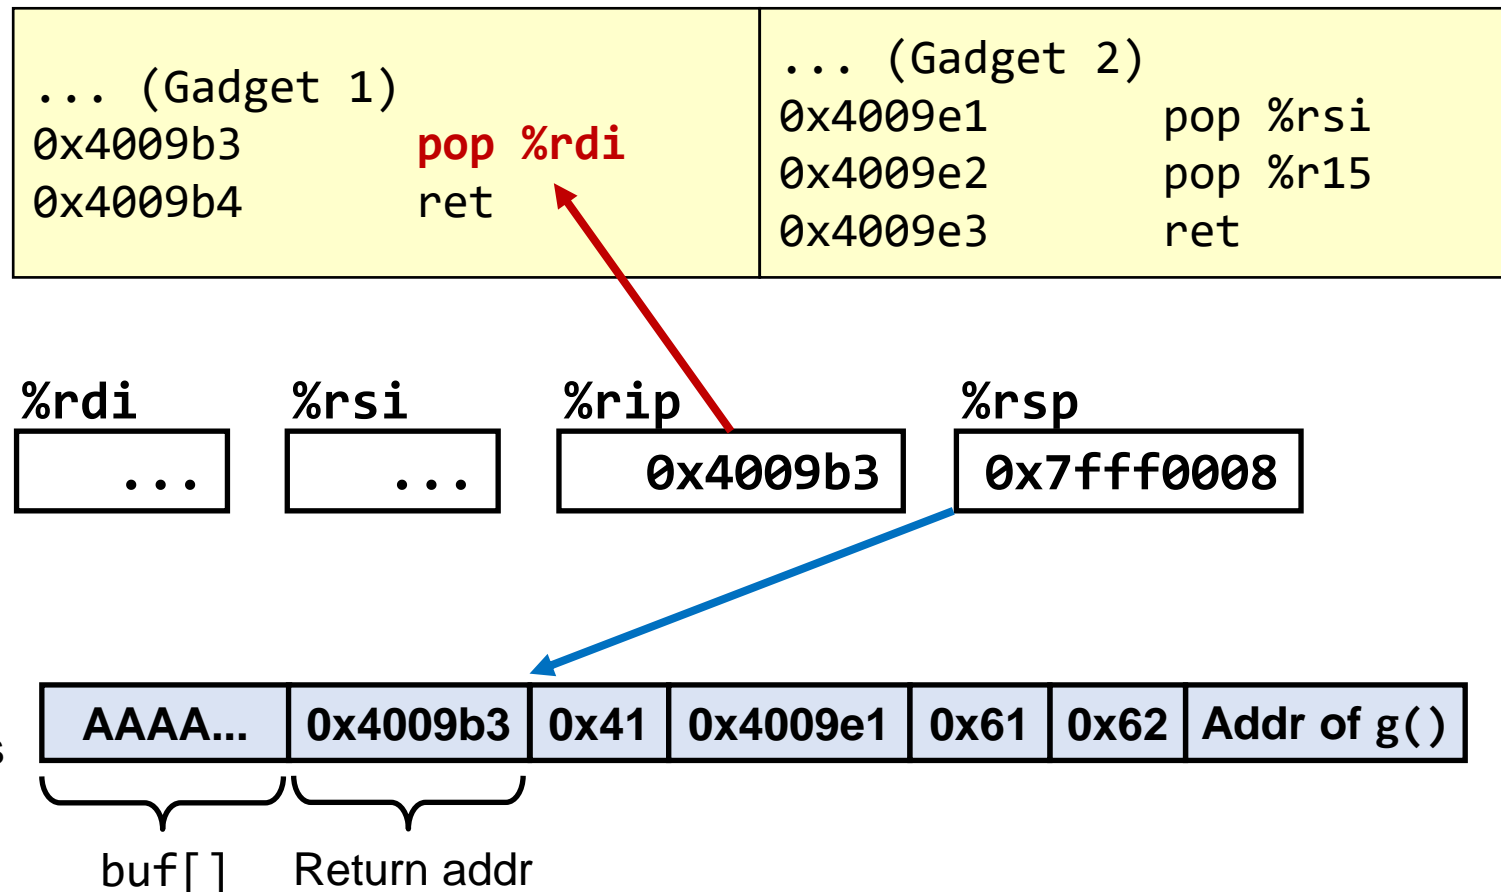
... (Gadget 1)	... (Gadget 2)
0x4009b3 pop %rdi	0x4009e1 pop %rsi
0x4009b4 ret	0x4009e2 pop %r15
	0x4009e3 ret

Quiz: What will be passed to function g()?



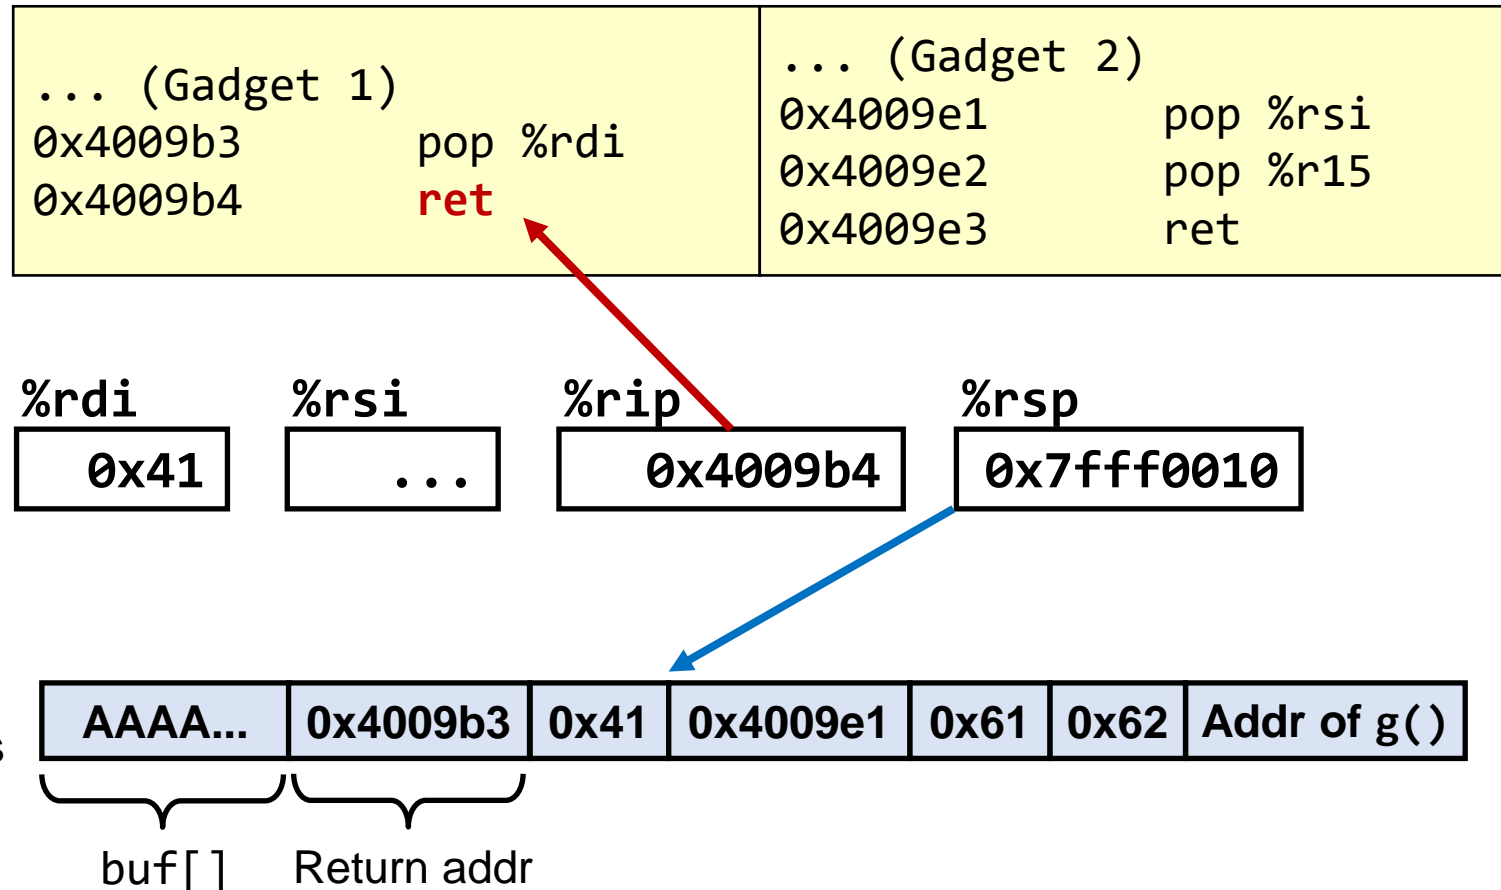
ROP Chain: Multiple Arguments

- The following pages explain this process step-by-step





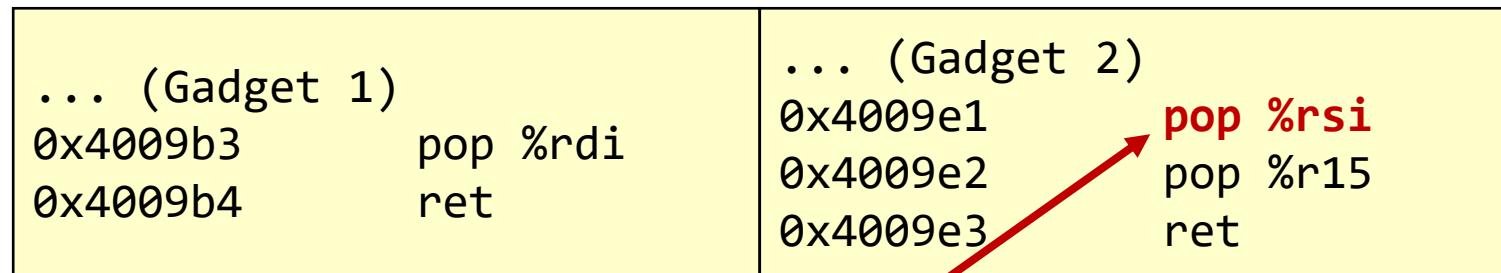
ROP Chain: Multiple Arguments

- The following pages explain this process step-by-step





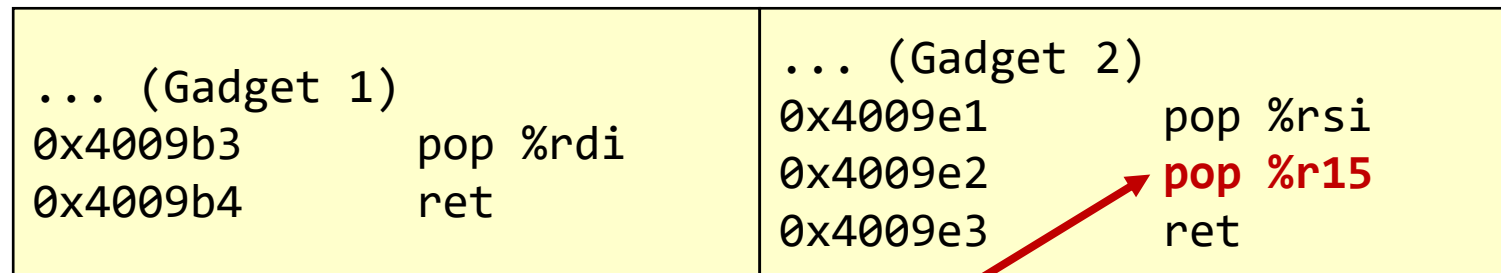
ROP Chain: Multiple Arguments

- The following pages explain this process step-by-step



ROP Chain: Multiple Arguments



- The following pages explain this process step-by-step



ROP Chain: Multiple Arguments

- The following pages explain this process step-by-step

... (Gadget 1)	... (Gadget 2)
0x4009b3 pop %rdi	0x4009e1 pop %rsi
0x4009b4 ret	0x4009e2 pop %r15
	0x4009e3 ret




ROP Chain: Multiple Functions

- Moreover, we can call a series of functions as well
- Simply put the next gadget right after the first function
 - In the example below, we place the **second 0x4009b3** immediately after the **address of f()**

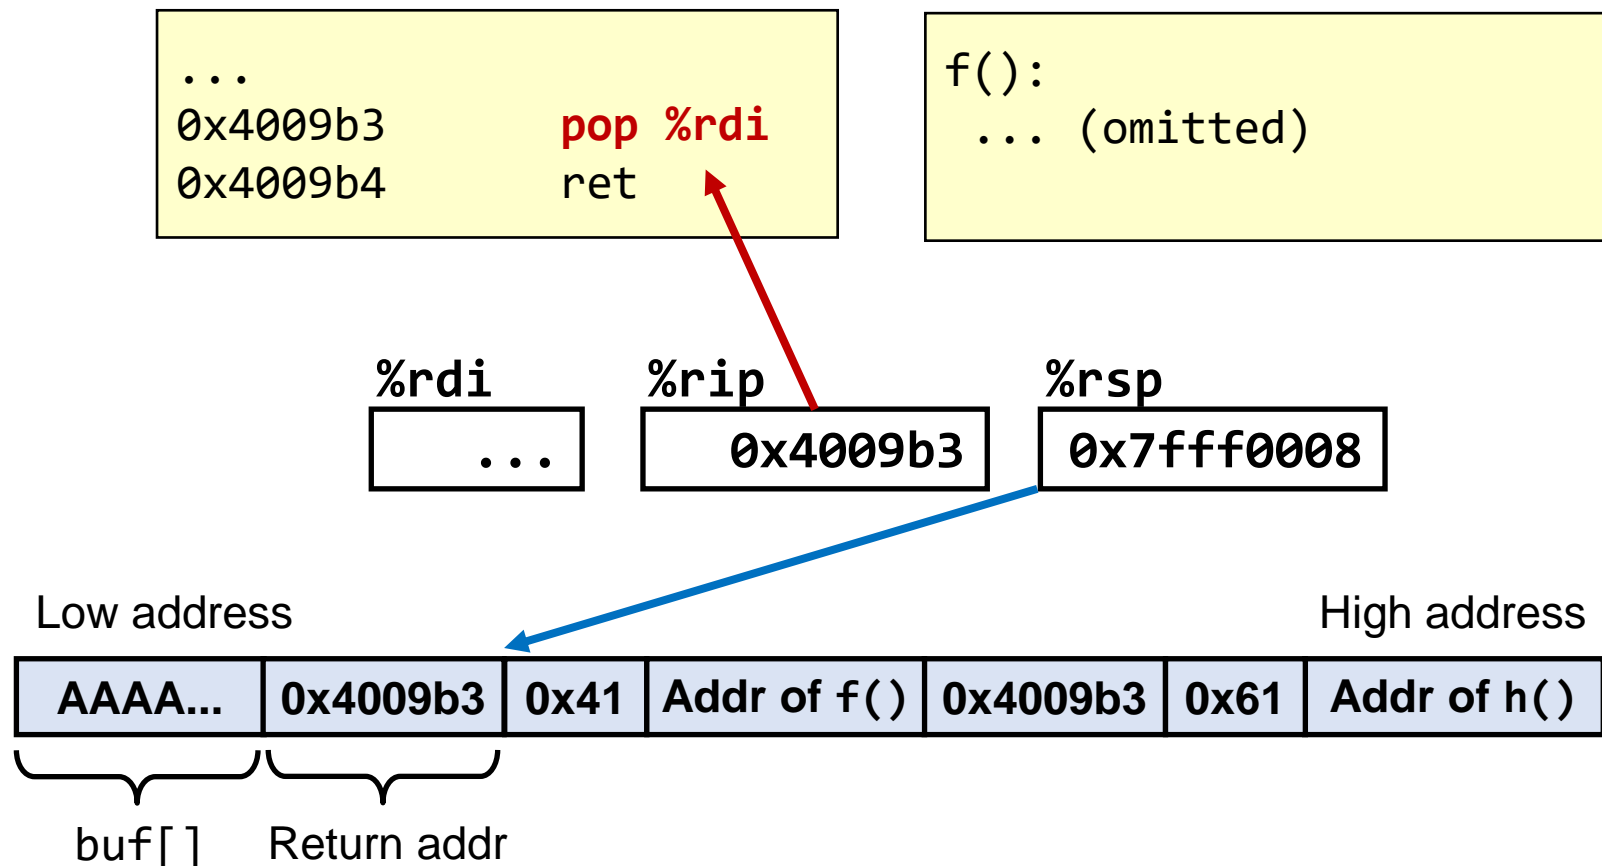
```
...  
0x4009b3      pop %rdi  
0x4009b4      ret
```

Quiz: Which functions will be called? What are their arguments?



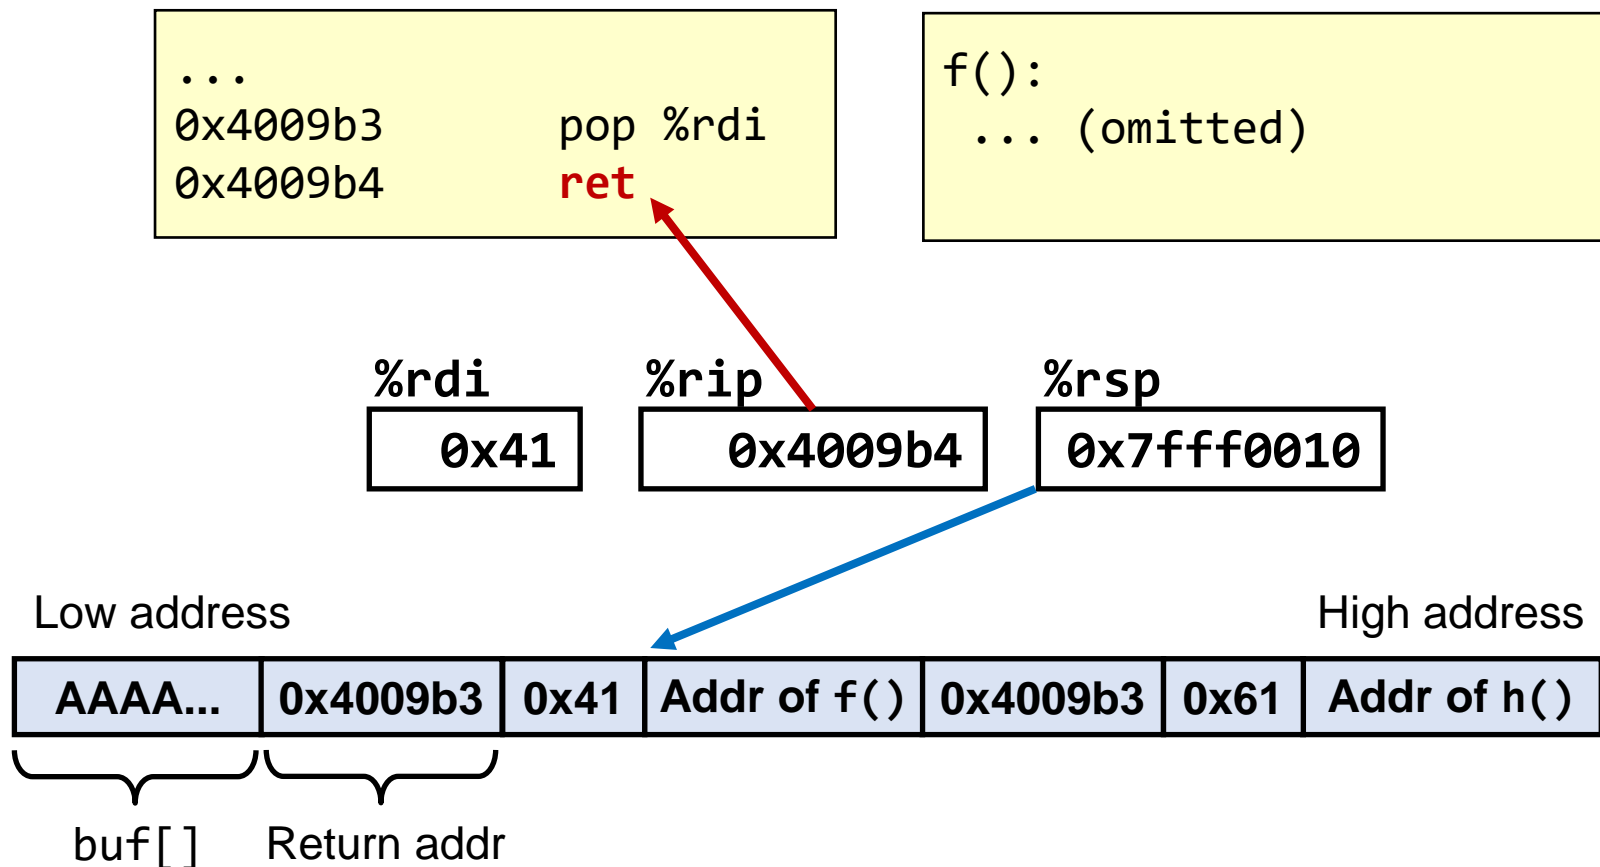
ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



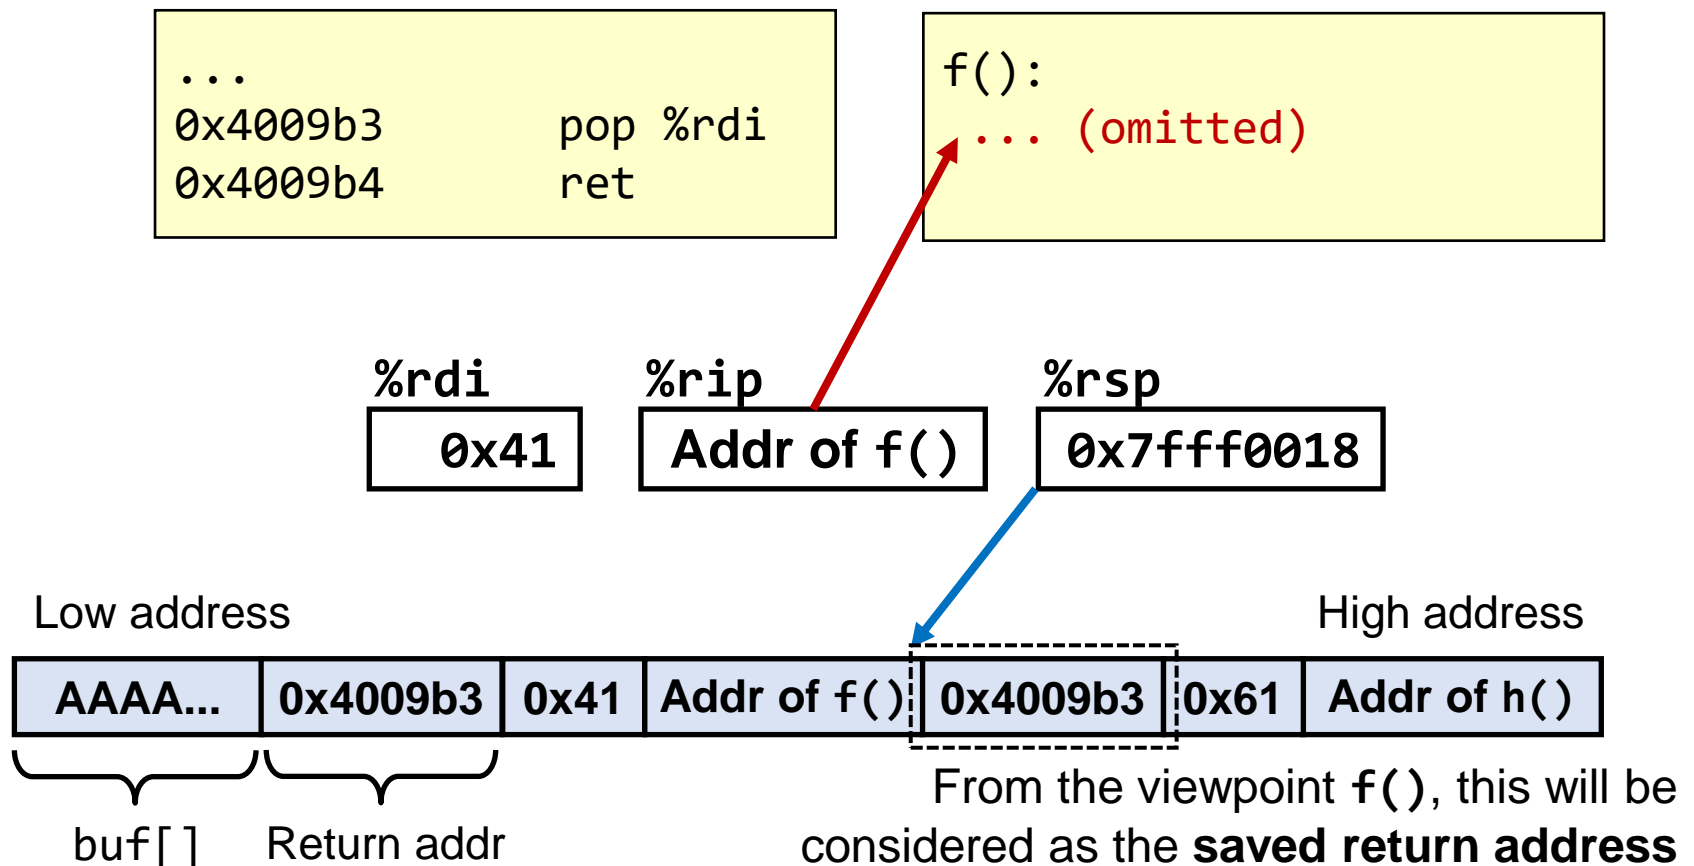
ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



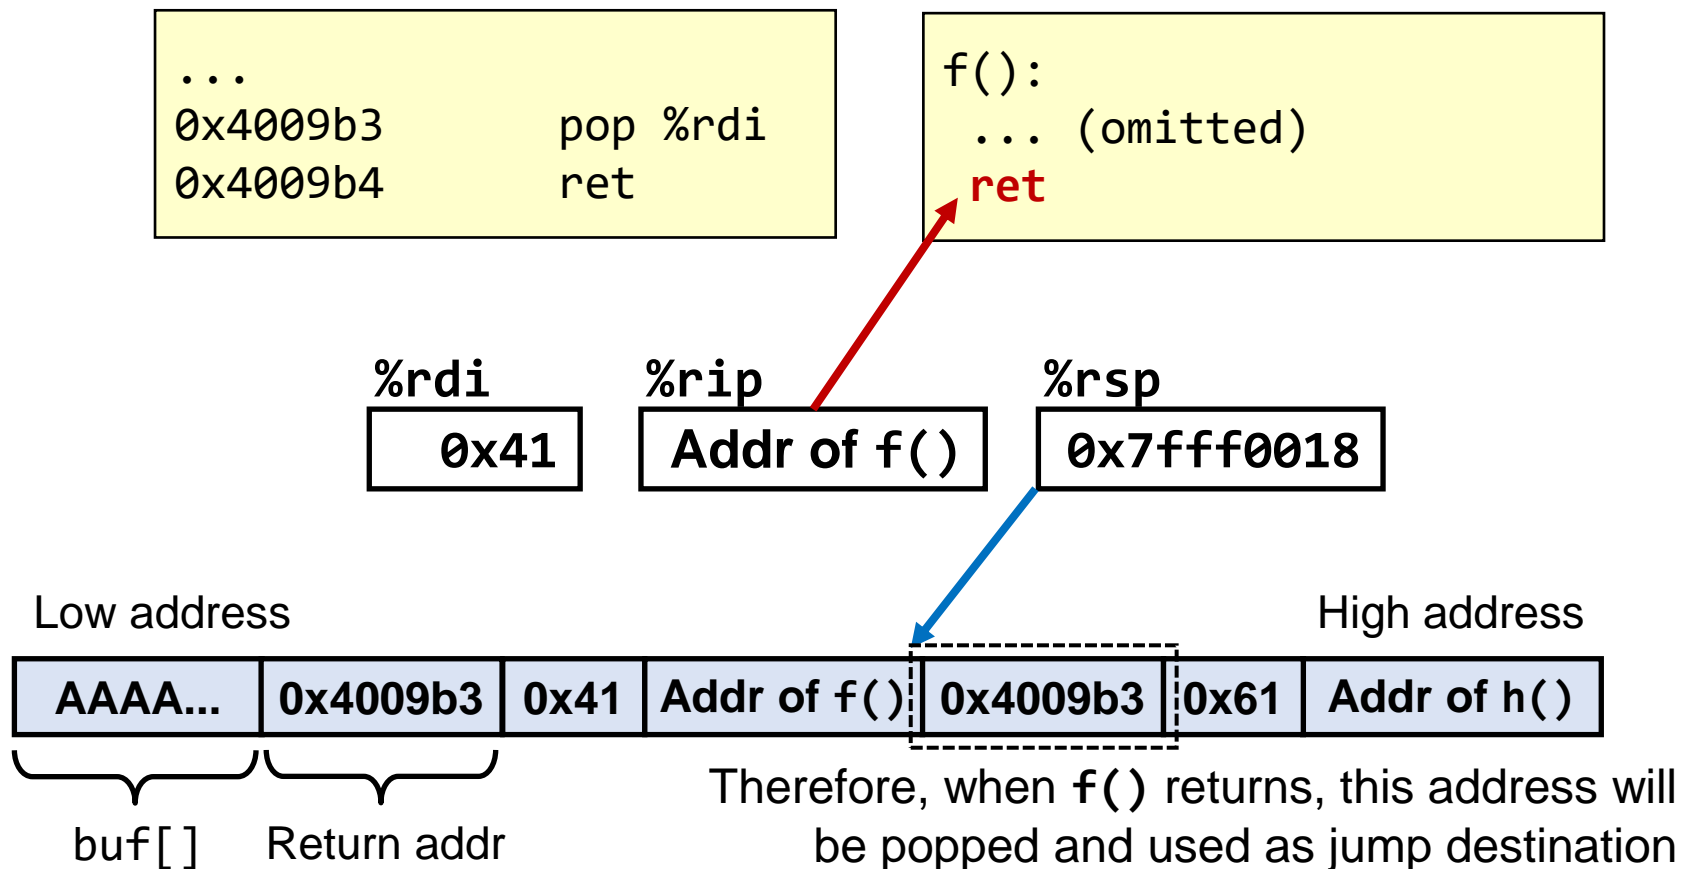
ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



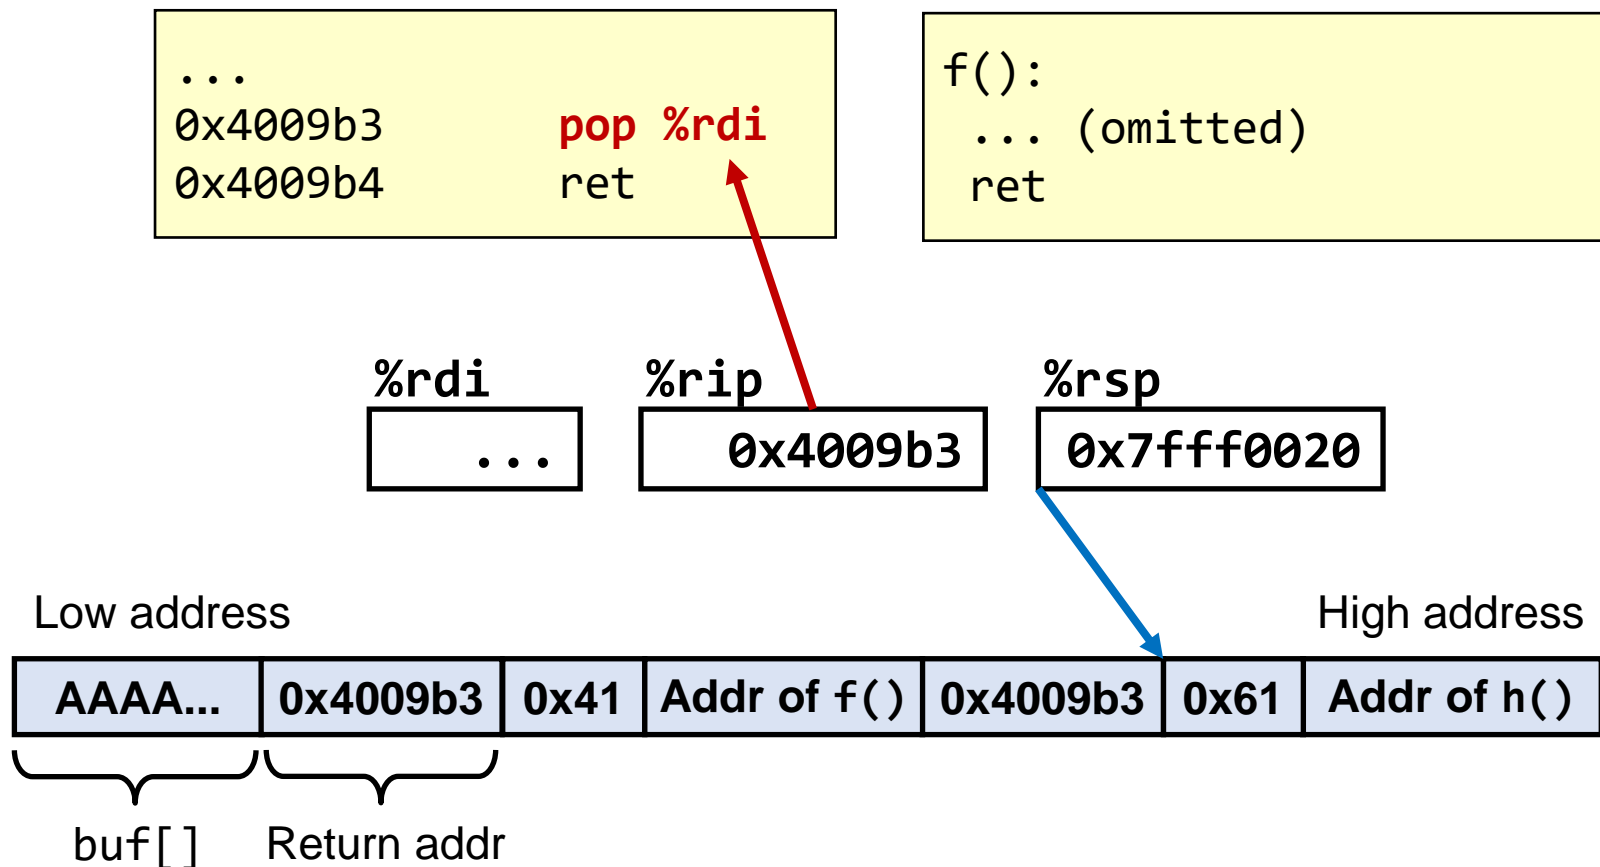
ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



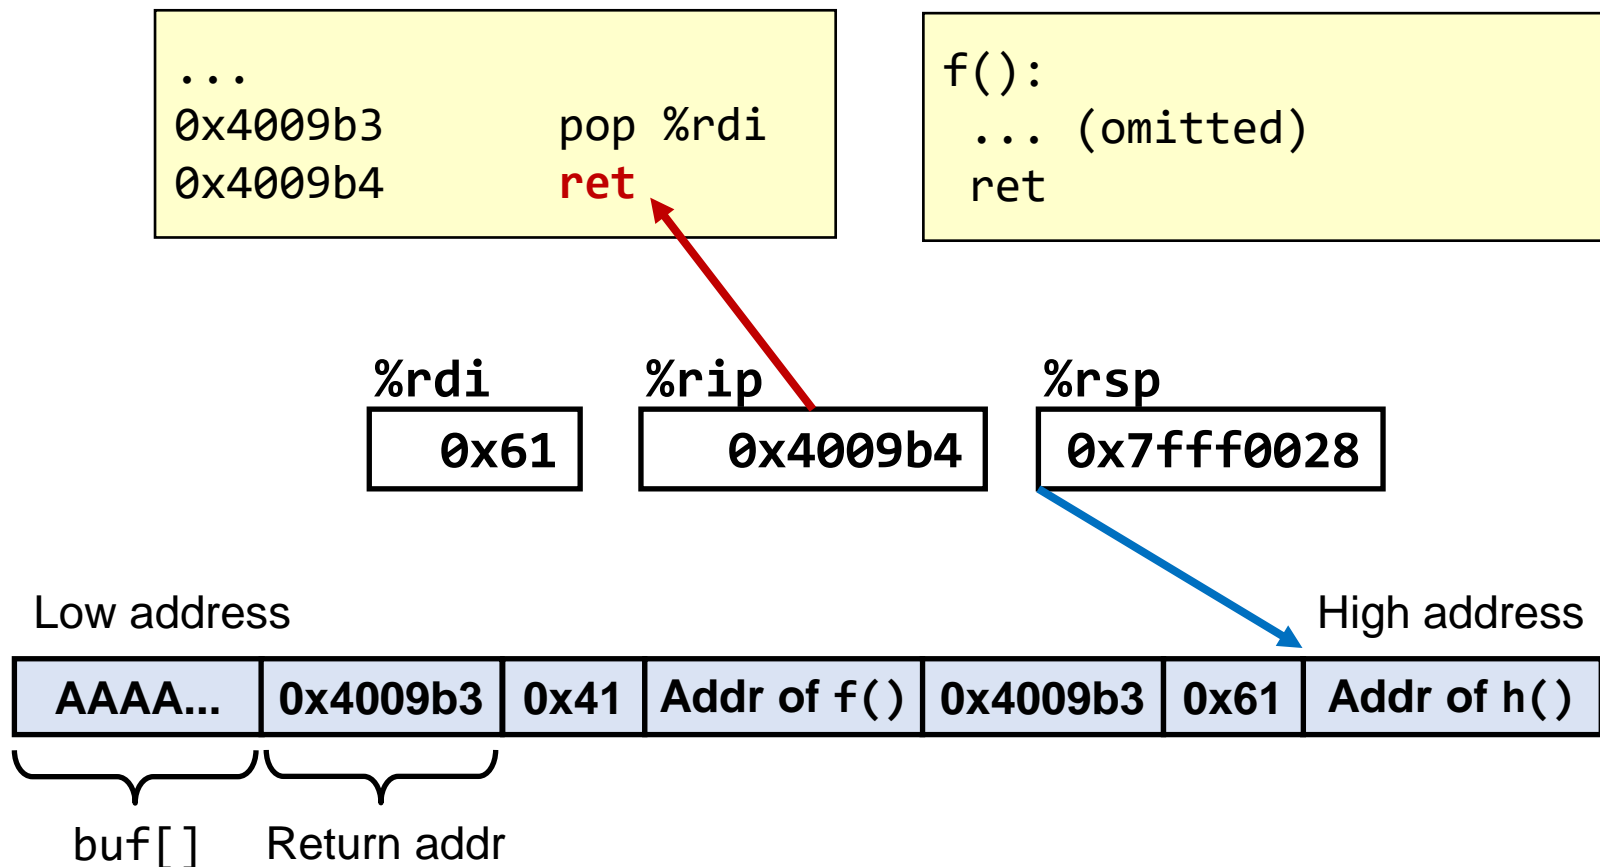
ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



ROP Chain: Multiple Functions

- The following pages explain this process step-by-step



Various Types of ROP Gadgets

- So far, we have discussed "**pop; ... ret;**" style gadgets
- But there can be other types of ROP gadgets, too
 - Ex) "add \$1000, %rcx; ret"
 - Ex) "xchg %rbx, %rdx; ret" # Exchange two registers
 - Ex) "mov %rax, (%rbx); ret" # Write to memory
- By chaining such gadgets, hacker can execute various logics and operations
 - This explains why it is called **return-oriented programming**


Using *ROP*, we have addressed
Challenge #1 (controlling arguments)

...

but what about ***Challenge #2 (ASLR)***?

Memory Disclosure Revisited

- Your program will contain many pointer variables that store addresses values
 - By disclosing (printing out) those values, hackers can get some clues about the memory layout



Disclosing Library Address

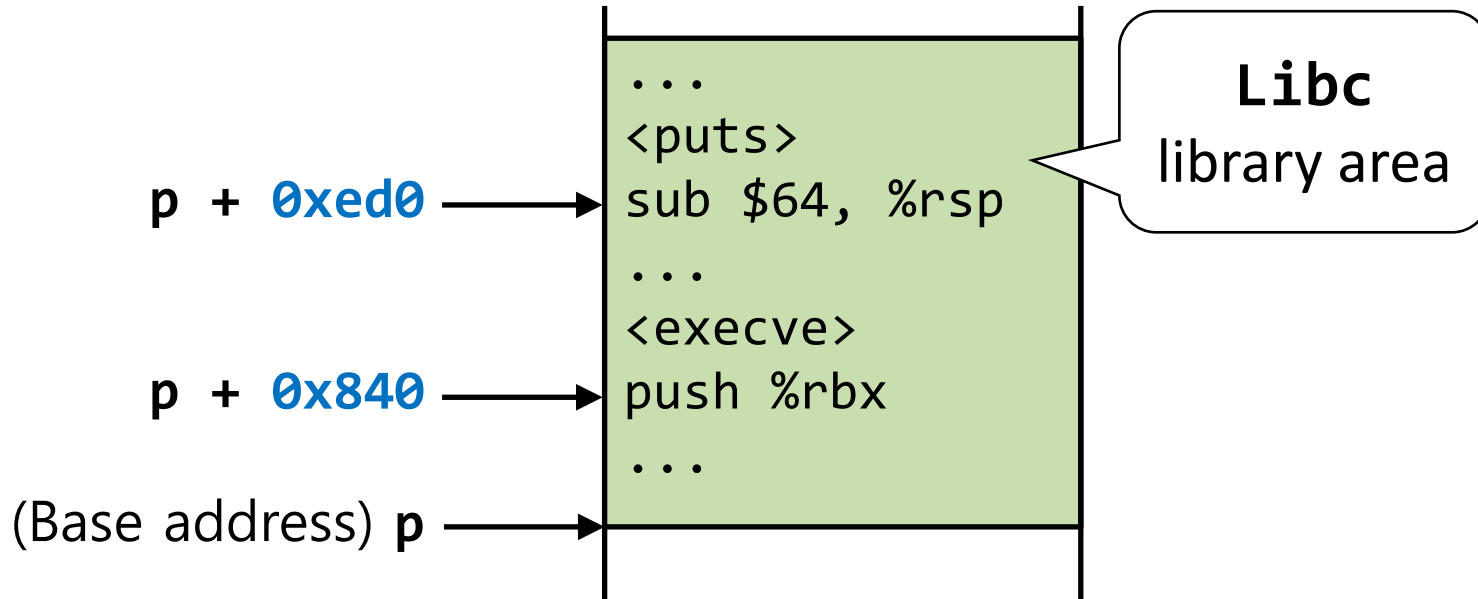
■ Let's consider the following code for example

- `dlopen()` + `dlsym()` are used to get the address of `puts()`
 - Don't have to know the details of these functions
- The obtained address will be stored in pointer `puts_fptr`
- Of course, this value will be different per execution (**ASLR**)

```
int main(void) {  
    int (*puts_fptr)(const char *);  
    void * handle = dlopen("libc.so.6", RTLD_LAZY);  
    puts_fptr = dlsym(handle, "puts");  
    puts_fptr("Hello world");  
    return 0;  
}
```

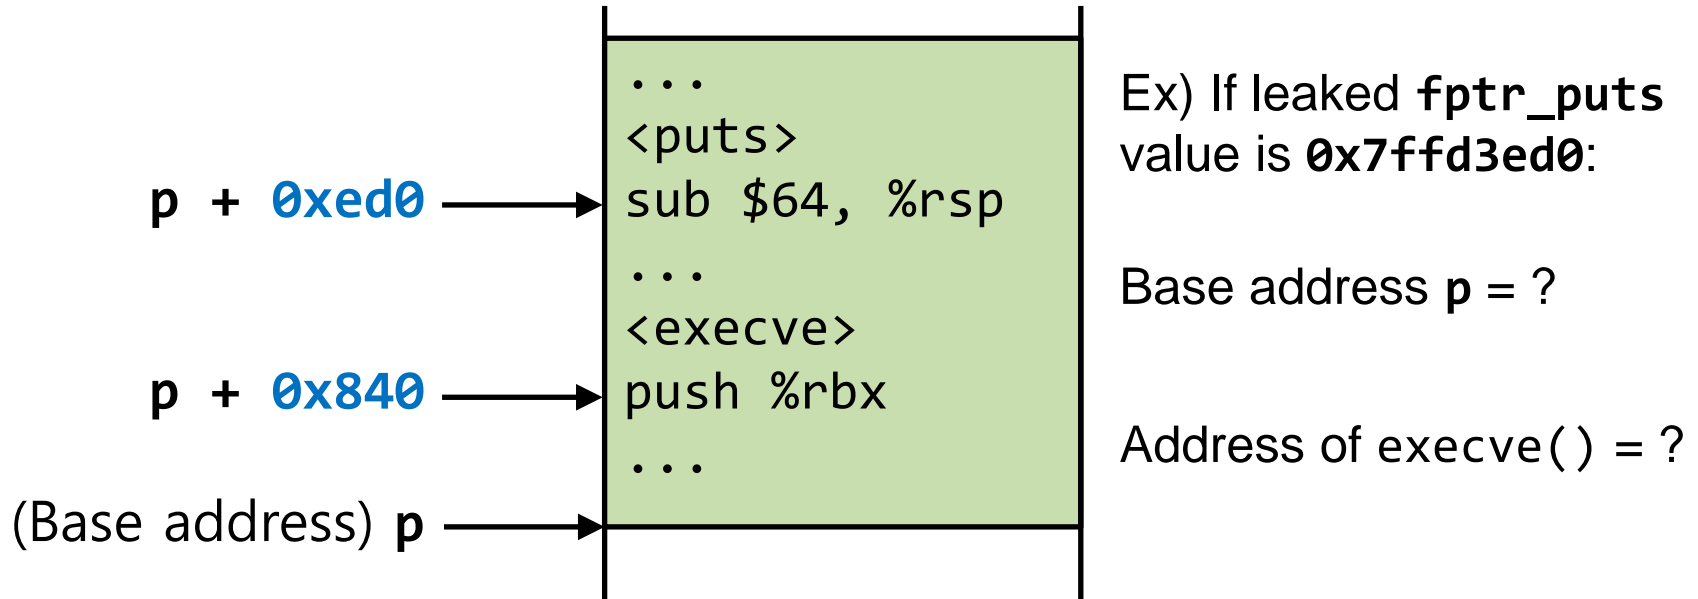
Disclosing Library Address

- Now, assume that hacker can disclose (print out) the **puts_fptr** variable in the previous example
 - For example, by using some buffer overflow (over-read)
- Then, the hacker figure out the address of library code
 - Why? The **offset** of a function within the library is fixed



Disclosing Library Address

- ASLR only randomizes the base address (the starting address) of the library in memory
 - Each function has unique and **fixed** offset (like `0xed0` for `puts`)
- Thus, by disclosing **`fptr_puts`**, hacker can also learn the base address (`p`) and the address of `execve()`



Does it really work?

- You may wonder if such scenario works in practice
- For example, let's consider a simple program below
 - The program has **no function pointer** to be leaked
 - Also, the BOF only allows memory corruption (**no disclosure**)
- Is this program still exploitable? (assuming ASLR)
 - **Yes, we can disclose the memory and bypass ASLR**

```
int main(void) {  
    char buf[8];  
    write(1, "Hello", 5);  
    read(0, buf, 160);  
    return 0;  
}
```


Background: Library Function Call

- To understand how it is possible, you should know what really happens during the library function call
- When a program calls a library function like `write()`, its `%rip` does not directly transfer to the library
 - Instead, it first moves to a **small code snippet called PLT**
 - This code snippet uses a **function pointer in a table called GOT**

```
int main(void) {  
    ...  
    write(1, "Hello", 5);  
    ...  
}
```

```
<main>  
...  
0x400579:  call 0x400430 <write@plt>  
...  
  
<write@plt>  
0x400430  jmp *0x601018 # GOT entry  
...
```

Background: PLT and GOT

- In other words, you can think that compiler and linker implicitly generate some function pointer table and fill it
 - To enable your program to call a function in library
 - Each library function called by your program has its PLT+GOT
 - GOT is filled at runtime (cannot be determined during compile)

PLT code snippets

```
<write@plt>
0x400430  jmp *0x601018
...

<read@plt>
0x400440  jmp *0x601020
...
```

GOT entries

```
0x601018: 0x7ffda8d0
0x601020: 0x7ffd2170
...
```

Library

```
...
<write>
...

<read>
...
```

So we can bypass ASLR by disclosing this GOT entry (function pointer)


How? By using ROP again

ROP for Memory Disclosure

- As before, we will overwrite the saved return address by using the buffer overflow in the example program

```
int main(void) {  
    char buf[8];  
    write(1, "Hello", 5);  
    read(0, buf, 160); // BOF  
    return 0;  
}
```


(Before the BOF. Assume each block is 8-byte)



ROP for Memory Disclosure

- By using ROP, we can call `write(1, &GOT, ...)`
 - And chain more gadgets to call more functions and **get a shell**

<main>	(Gadgets)
...	...
0x400579: call 0x400430 <write@plt>	0x4009b3 pop %rdi
...	0x4009b4 ret
	...
<write@plt>	0x4009c5 pop %rsi
0x400430 jmp *0x601018 # GOT entry	0x4009c6 ret
...	



Notes on Memory Disclosure

■ Is library function offset always predictable?


- Depending on your Linux version, `Libc` library version will vary and the offset of each function will change, too
- Even without an access to the `Libc` file, attackers can still infer the version of `Libc` (we will not discuss this deeply)

■ The idea of figuring out library address can be applied to other memory areas, too

- For example, the address of stack or heap can be also inferred

Position-independent Executable

- Nowadays, ASLR is applied to Code & Data sections too
 - This is called position-independent executable (PIE)
- For this, compiler must generate complex assembly code
 - So we will not cover it in this course



Lessons

- **Hackers are more persistent than you think**
 - They often come up with creative methods to bypass mitigation
- **So the impact of software vulnerabilities should not be underestimated or overlooked**
- **To precisely understand the outcome of a bug, it is important to know the internals of computer systems**
 - Ex) If you didn't know the existence of PLT/GOT, it would be hard to imagine how the exploit is possible