

Chapter 7. Heap Vulnerability

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering

Sogang University

Topics

■ Brief background of dynamic memory allocation*

- Linked list of free blocks
- Heap metadata

■ Buffer overflow in heap

- Common patterns of heap buffer overflow
- Exploiting heap buffer overflow

■ Use-after-free

- Principle and examples of use-after-free
- Exploitation of use-after-free vulnerability

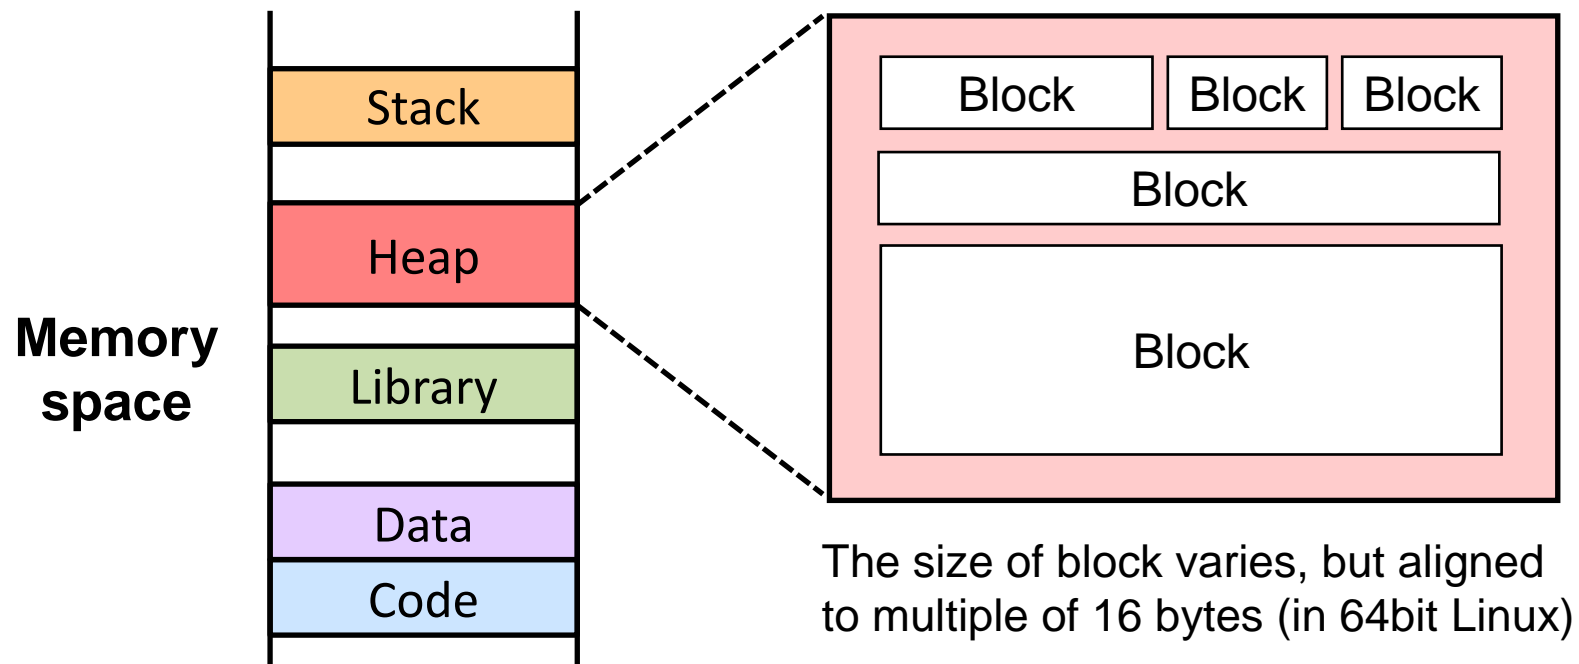
Dynamic Memory Allocation

- **Often, the size of data to process is decided at runtime**
 - Ex) Length of input string is provided by user
 - Ex) Must add a node to the linked list whenever a user requests
- **We can use dynamic memory allocation**
 - `void *malloc(size_t size);`
 - `void free(void *ptr);`

```
// Size of buffer is decided by user input
size_t input_size = read_size_t();
char *p = (char*) malloc(input_size);
if (p == NULL)
    return;
fgets(p, input_size, stdin);
```

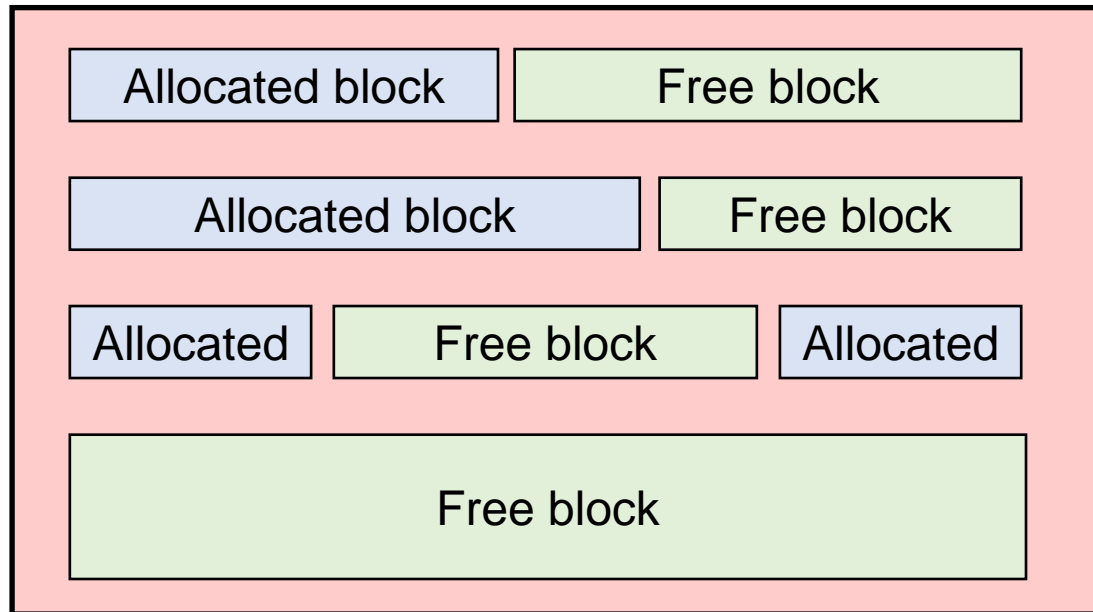
Layout of Heap Memory

- Memory allocator can increase the size of heap area **(the red section)** by invoking system calls like `sbrk()`
- Inside the heap memory area, small memory blocks (sometimes called chunks) are created and managed



Heap Memory Block

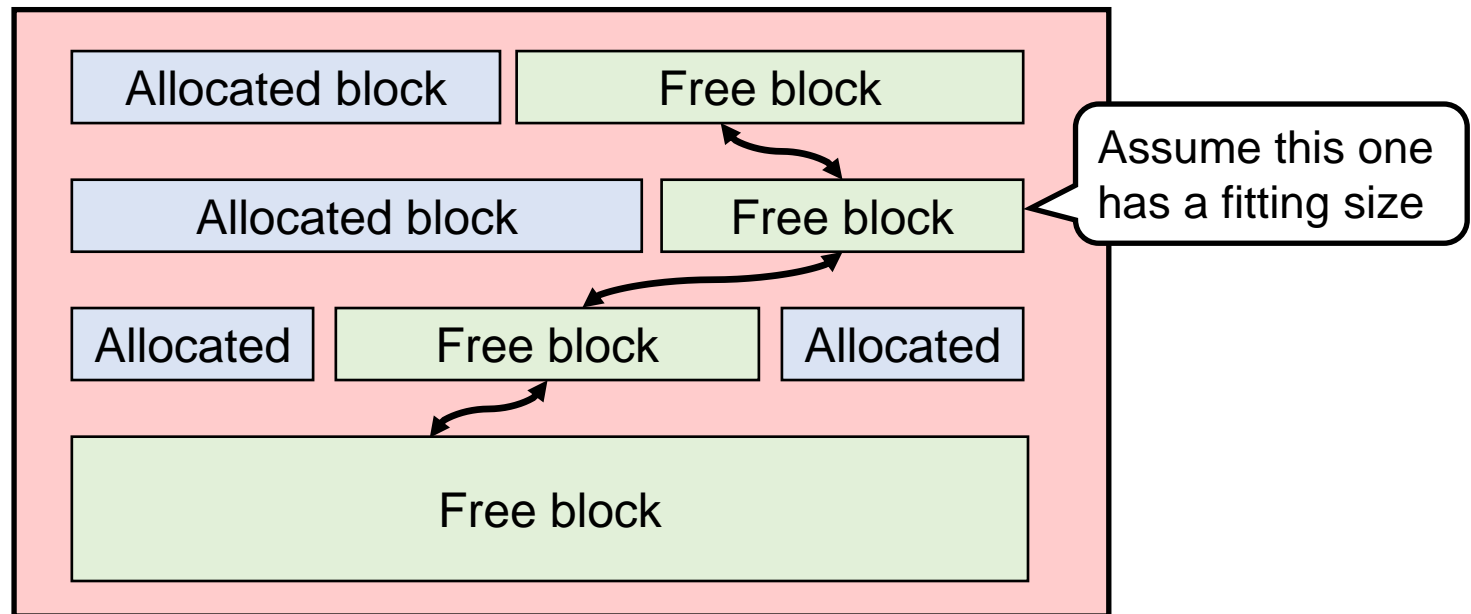
- When the program calls `malloc()`, the memory allocator will return one of the available blocks
 - Block that is returned by `malloc()` is called ***allocated*** block
 - Other blocks are called ***free (not allocated)*** blocks



List of Free Blocks

■ The memory allocator manages linked list of free blocks

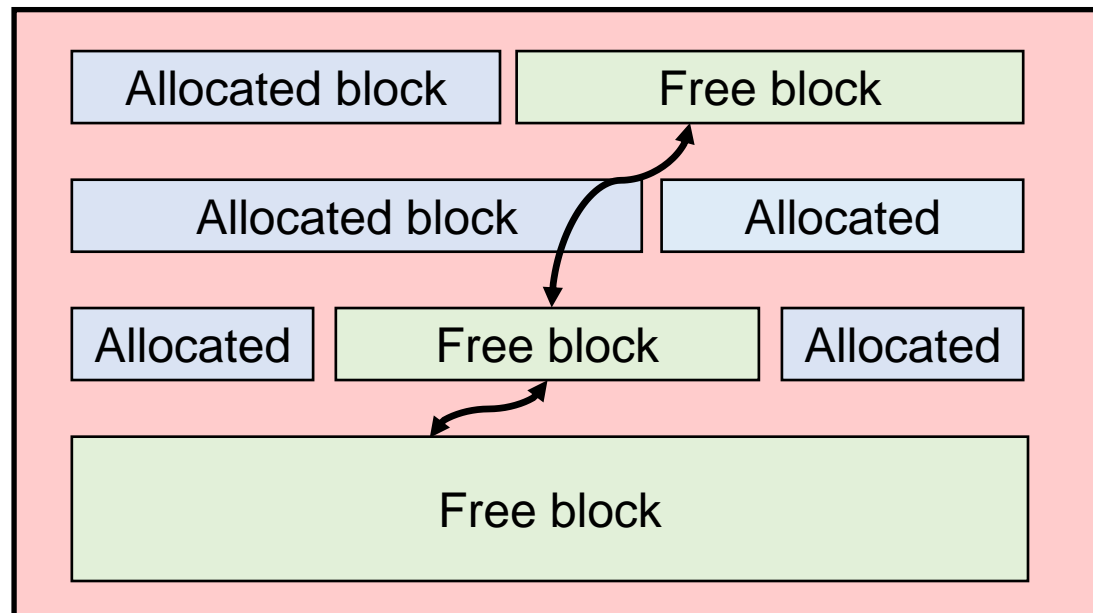
- When the program calls **p = malloc()**:
 - A block of fitting size is found from the list and returned
 - Or a free block with larger size can be split and returned



List of Free Blocks

■ The memory allocator manages linked list of free blocks

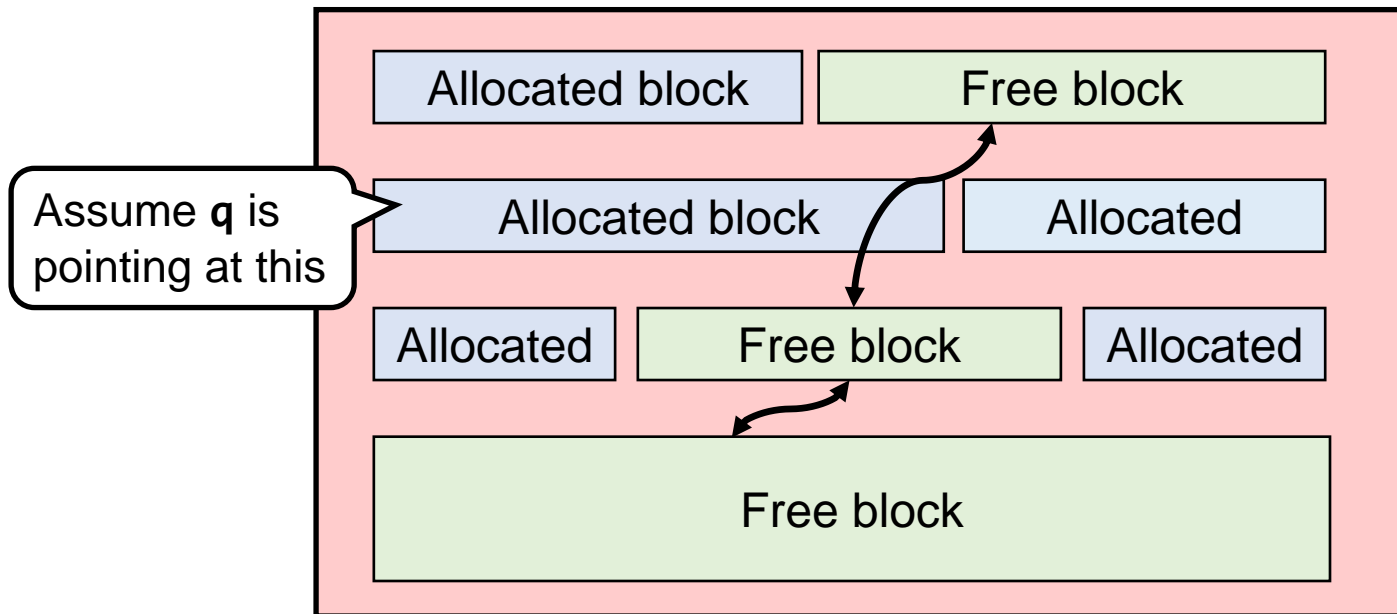
- When the program calls **p = malloc()**:
 - A block of fitting size is found from the list and returned
 - Or a free block with larger size can be split and returned



List of Free Blocks

■ The memory allocator manages linked list of free blocks

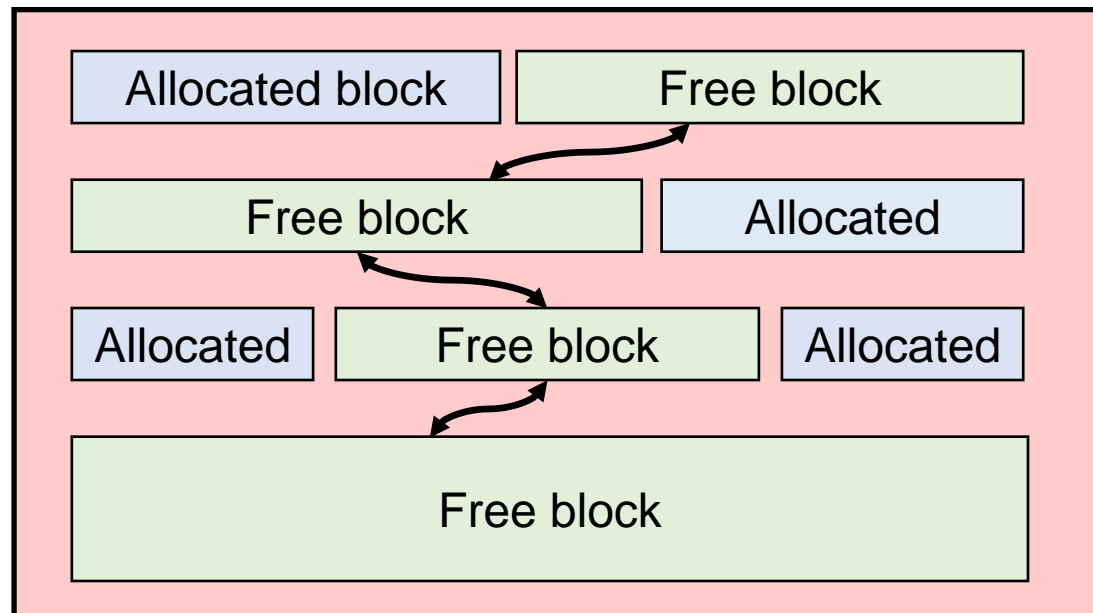
- When the program calls **free(q)**:
 - The block pointed by **q** is inserted to the list of free blocks
 - Now the block space is used to store a pointer to next node



List of Free Blocks


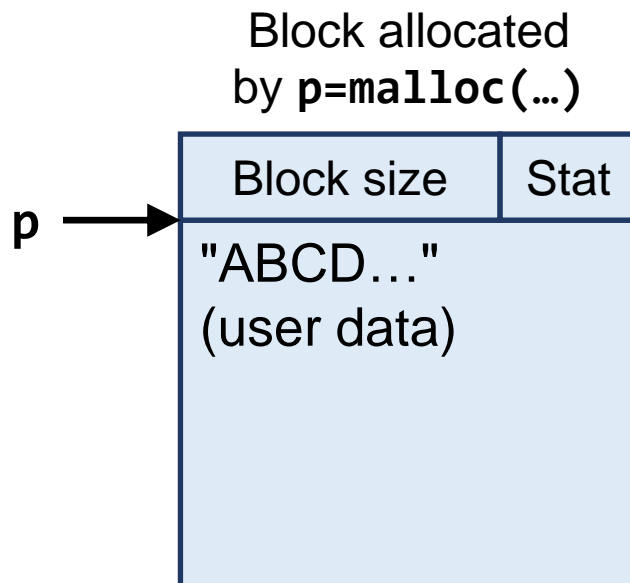
■ The memory allocator manages linked list of free blocks

- When the program calls **free(q)**:
 - The block pointed by **q** is inserted to the list of free blocks
 - Now the block space is used to store a pointer to next node



Metadata in Heap Memory Block

- To manage blocks in this way, the memory allocator must store various information other than user data
 - Ex) Size of block, status of current/adjacent block, pointer to next/previous node of a doubly linked list ...
 - Such additional information is often called *metadata*



Real-world memory allocators are much more complex than this

But this will be enough for our course

Topics

■ Brief background of dynamic memory allocation

- Linked list of free blocks
- Heap metadata

■ Buffer overflow in heap

- Common patterns of heap buffer overflow
- Exploiting heap buffer overflow

■ Use-after-free

- Principle and examples of use-after-free
- Exploitation of use-after-free vulnerability

Buffer Overflow in Heap

- **Heap BOF can occur for similar reasons to stack BOF**
 - Ex) Calling unsafe library functions like `gets()` or `scanf("%s")`
 - Ex) Allowing array to be accessed with arbitrary index
- **Moreover, it is easier to make mistakes in heap because the allocation size is usually affected by user input**
 - Consider the example below: what can go wrong here?

```
uint item_count = read_uint();
int *arr = (int*) malloc(item_count * 4); // int is 4-byte
if (arr == NULL)
    return;
uint idx = read_uint();
if (idx < item_count)
    arr[idx] = 1;
```

Buffer Overflow in Heap

- **Heap BOF can occur for similar reasons to stack BOF**
 - Ex) Calling unsafe library functions like `gets()` or `scanf("%s")`
 - Ex) Allowing array to be accessed with arbitrary index
- **Moreover, it is easier to make mistakes in heap because the allocation size is usually affected by user input**
 - Consider the example below: what can go wrong here?


```
uint item_count = read_uint();  
int *arr = (int*) malloc(item_count * 4); // int is 4-byte  
if (arr == NULL)  
    return;  
uint idx = read_uint();  
if (idx < item_count)  
    arr[idx] = 1;
```

Input: 0x40000001

Size: 0x4 (*integer overflow*)

Exploiting Heap BOF

- **How can an attacker exploit BOF that occurs in heap?**
 - Unlike stack, there is **no saved return address** in heap
- **First, corruption of a structure (or object in C++) may lead to security issues, such as control hijack**
- **Second, the corruption of metadata in heap blocks can cause problems as well**



Corruption of Structure/Object

- Assume a structure (or object in C++) allocated in heap
- If there is a field (or property) whose type is **function pointer**, corrupting such field can lead to control hijack
 - C++ object often contains a pointer to function pointer table (which is known as **virtual table**)

```
struct S {  
    char buf[16];  
    void (*handler)(char *s);  
};  
  
void f(void) {  
    struct S *s = malloc(sizeof(struct S));  
    gets(s->buf); // Buffer overflow (may corrupt "handler")  
    s->handler("input msg");  
}
```


Corruption of Metadata

- Recall that memory allocator often maintains a (doubly) linked list of free blocks
- What if the Next or Prev pointer within a free block is corrupted by buffer overflow?
 - When such block is removed from the linked list, dangerous memory operations can occur

Block size	Stat
Next (corrupted)	
Prev (corrupted)	
Block size	

```
remove_from_list(b) {  
    Block *n = b->next;  
    Block *p = b->prev;  
    n->prev = p;  
    p->next = n;  
}
```

Modern Memory Allocators

- Recent memory allocators are equipped with many protection mechanisms against metadata corruption
 - The code below shows a patch introduced in 2004
- As a result, nowadays it is rather hard to exploit a heap buffer overflow by corrupting block metadata

```
#define unlink(P, BK, FD) {  
    FD = P->fd ;  
    BK = P->bk ;  
    if (FD->bk != P || BK->fd != P) error(); // Validation  
    else {  
        FD->bk = BK;  
        BK->fd = FD;  
    }  
}
```

Topics

■ Brief background of dynamic memory allocation

- Linked list of free blocks
- Heap metadata

■ Buffer overflow in heap


- Common patterns of heap buffer overflow
- Exploiting heap buffer overflow

■ Use-after-free

- Principle and examples of use-after-free
- Exploitation of use-after-free vulnerability


Use-After-Free (UAF)

- Another popular source of heap vulnerability
- Accessing a heap memory location that has already been deallocated with `free()`
 - Has been a well-known concept in computer science
 - Sometimes called *use of a dangling pointer*
 - Gained lots of interest from security researchers around 2010



Use-After-Free (UAF)

- Another popular source of heap vulnerability
- Accessing a heap memory location that has already been deallocated with `free()`
 - Has been a well-known concept in computer science
 - Sometimes called *use of a dangling pointer*
 - Gained lots of interest from security researchers around 2010



Use-After-Free Example

- In the code below, the heap memory block pointed by `struct S *s` is deallocated
- But this memory location is accessed after `free()`
 - After the deallocation, cannot guarantee what is stored there

```
struct S {  
    int x;  
    ...  
};  
  
struct S *s = (struct S*) malloc(sizeof(struct S));  
...  
free(s);  
...  
printf("x = %d\n", s->x); // Accessing deallocated memory
```

More Realistic Example (1)

- Assume that there are multiple pointers that are pointing at the allocated memory block
- It can be hard to remember all the pointer variables that are related to the deallocated memory block

```
struct S *s1 = (struct S*) malloc(sizeof(struct S));
struct S *s2 = s1;
px = &(s1->x);
...
free(s1); // OK, let's not use 's1' anymore
...
// But forgot about 's2' or 'px'
s2->x = 100;
printf("x = %d\n", *px);
```

More Realistic Example (2)

- Assume a global pointer variable pointing at heap block
- If the pointed heap block is deallocated, must invalidate the pointer (for example, by setting into NULL)
 - Forgetting to do so may lead to *use-after-free*

```
struct S* items[32]; // Each element is a pointer to struct

void delete(int idx) {
    if (idx < 0 || idx >= 32)
        return;
    if (items[idx] != NULL) {
        free(items[idx]);
    }
    // Didn't we forget something here?
}
```


Security Impact of Use-After-Free

- Now let's think about the severity of this bug
- Consider the example code below
 - `printf()` after `free()` will print out old data in the free block
 - Following `s->x=100;` will update unused space in the free block
 - In some case, these may disclose or corrupt the heap metadata
- But is that all?
 - Or is it possible to exploit UAF and **hijack the control-flow?**

```
struct S *s = (struct S*) malloc(sizeof(struct S));  
...  
free(s);  
printf("x = %d\n", s->x); // Using 's'  
s->x = 100;
```

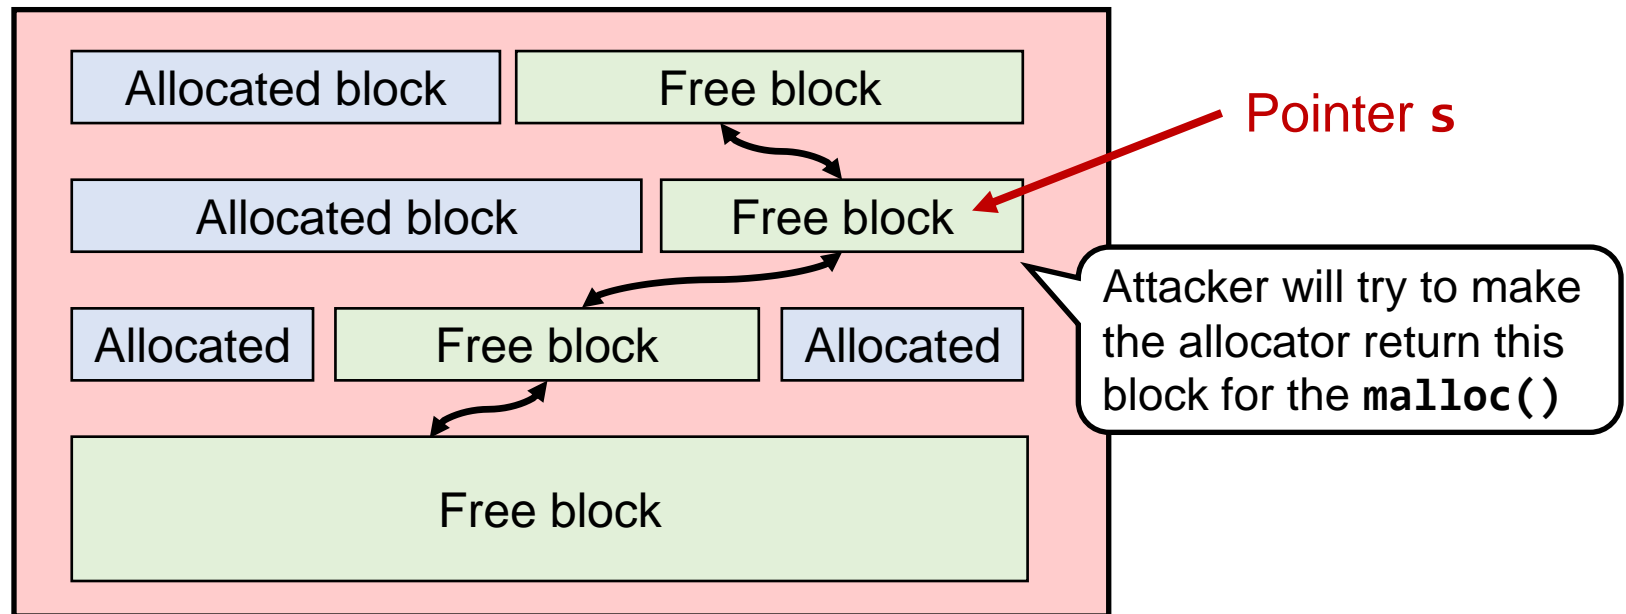
Free blocks can be re-allocated

- In practice, often there is an interval between the `free()` and use-after-free
- What if there are `malloc()` calls within that interval?
- Memory block pointed by 's' may be allocated again

```
struct S *s = (struct S*) malloc(sizeof(struct S));
...
free(s);
...
malloc(...);
malloc(...);
...
printf("x = %d\n", s->x); // Using 's'
s->x = 100;
```

Controlling the Heap Memory


- Recall that free blocks are managed as a linked list
 - The allocator will find a block with *fitting size* from this list
- By carefully deciding the size of `malloc()`, the attacker can enforce the wanted block to be returned



Exploiting Use-After-Free

- Assume that struct **Data** has a character array field (**buf**) and a function pointer field (**handler**)

```
➔ struct Data *p = (struct Data*) malloc(sizeof(struct Data));  
...  
free(p);  
...  
len = read_size_t();  
char *str = (char *) malloc(len); // Will read in string  
...  
p->handler(p->buf); // Use-after-free
```



1. struct Data allocated

Exploiting Use-After-Free

- Assume that struct **Data** has a character array field (**buf**) and a function pointer field (**handler**)

```
struct Data *p = (struct Data*) malloc(sizeof(struct Data));  
...  
→ free(p);  
...  
len = read_size_t();  
char *str = (char *) malloc(len); // Will read in string  
...  
p->handler(p->buf); // Use-after-free
```




2. Deallocated with `free()`

Exploiting Use-After-Free

- Assume that struct **Data** has a character array field (**buf**) and a function pointer field (**handler**)

```
struct Data *p = (struct Data*) malloc(sizeof(struct Data));  
...  
free(p);  
...  
len = read_size_t();  
→ char *str = (char *) malloc(len); // Will read in string  
...  
p->handler(p->buf); // Use-after-free
```




3. Allocated again as `char[]`

Exploiting Use-After-Free

- Assume that struct **Data** has a character array field (**buf**) and a function pointer field (**handler**)

```
struct Data *p = (struct Data*) malloc(sizeof(struct Data));  
...  
free(p);  
...  
len = read_size_t();  
char *str = (char *) malloc(len); // Will read in string  
...  
→ p->handler(p->buf); // Use-after-free
```



4. Interpreting `char[]`
as `struct Data`

Control Hijack!

Use-After-Free in Stack?

- Can use-after-free occur in stack as well?
- There is no `malloc()` in stack, but the allocation and deallocation of stack frame is conceptually similar
- If a function **returns a pointer to its local variable**, use of that pointer will access invalid memory location
 - But this type of bug occurs rarely (easy to avoid)

```
char *f(void) {  
    char buf[32];  
    char *p = buf;  
    return p; // But buf does not exist after f() returns  
}
```


Side-note: Uninitialized Data Use

- UAF can be thought as *accessing memory too late*
- In contrast, program may also access *memory too early*
 - If a variable is used without being initialized, garbage values stored there will be used (use-before-initialize)
 - But the term ***uninitialized data/variable use*** is more popular
 - Can occur in any memory region (*stack, heap, or data*)
- Memory disclosure or even code execution can be possible in some cases

```
void f(void) {  
    struct S s;  
    read(0, s.buf, 32); // Using uninitialized pointer value  
    ...  
}
```