

Chapter 9. Kernel Security

Prof. Jaeseung Choi

Dept. of Computer Science and Engineering

Sogang University

Topics

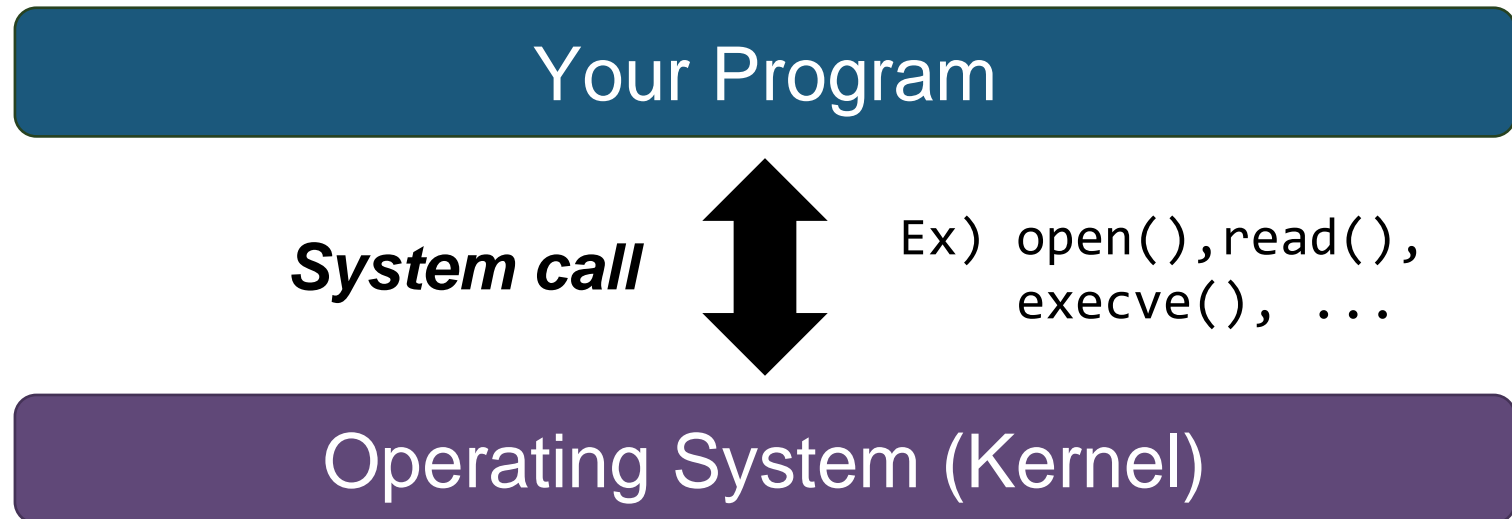
- **Software vulnerabilities in OS kernel**
 - Using user-provided pointer without check
 - Double fetch
 - Null dereference
 - Memory disclosure due to alignment

From User Code to Kernel Code

- So far, we've discussed various software vulnerabilities
- We have assumed **user code** during this discussion
 - Usually, what you write and run is user-level code
- Now let's take a look into the **kernel code**
 - The core part of operating system
 - Kernel code is executed in privileged mode
 - In **PintOS** project, you have written (or will write) kernel code
 - What kind of security issues can rise here?

Review: System Call

- Recall that the program code you wrote is not allowed to access the system resources directly
 - Ex) Reading a file or writing to a file
- To do that, you must make a request via **system calls**
 - Then, the OS will do the task for you



Review: `syscall` Instruction


- At assembly level, we must use `syscall` instruction
- Similar to function call, but the function is in kernel
- When the code sets particular registers properly and executes this instruction, a system call is invoked
 - Now, let's dig a little bit deeper about the system call

```
...  
mov    $0, %rsi      # %rsi must contain flag (option)  
mov    ..., %rdi     # %rdi must point to filename string  
mov    $0x2, %rax    # System call ID of open() is 2  
syscall
```

What happens at this moment?

■ Upon the execution of `syscall` instruction:

- The CPU mode changes to privileged mode (kernel mode)
- Kernel-level memory space becomes accessible
- Program counter moves to pre-defined address in kernel code



System Call Handler in Kernel

- Now it's the turn of **system call handler** in kernel code
 - It examines the system call arguments passed from the user
 - Ex) **open** system call: What is the name of file that the user requests to open? Does this user have a proper permission?
 - Ex) **read** system call: What is the input file descriptor? Where does the user want the file content to be stored?
- **Note: To do these things, system call handler code will have to access both the user space and kernel space**

Vulnerabilities in Kernel

- Most of the software **vulnerabilities that we learned** until now can also occur in the kernel code
- For example, buffer overflow can occur and corrupt the critical data stored in the kernel memory
 - Ex) `strcpy()` below copies string without any length check
- **Kernel must not trust user-provided arguments**
 - An attacker may invoke malicious system calls and exploit kernel vulnerabilities, in order to obtain the **root** privilege

```
// User-provided arguments 'filename' and 'flags'
int open_handler(char *filename, int flags) {
    char buf[32];
    strcpy(buf, filename); // Buffer overflow
    ...
}
```



OK, then is this all? (Buffer overflow, format string bug, use-after-free, ...)

No, there are other **unique kind of bugs that can occur in the kernel code**

Danger in User-Provided Pointer

- Let's consider the system call handler for `read()`
 - Assume that the file content is first loaded into kernel space
 - In the end, this content must be copied to the user's buffer (`buf`)
 - In normal case, `buf` must be an address in user space
- But what if the provided `buf` is a kernel-space address?
 - `memcpy()` may overwrite and corrupt kernel memory

```
read_handler(int fd, void *buf, size_t n) {  
    ... // Read in the file content  
    memcpy(buf, file_content, n);  
    return 0;  
}
```



Review: Pintos Manual

- In fact, if you have taken the *Operating System* course, you must be already familiar with all these issues
- The Pintos manual also explains this vulnerability
 - Section 3.1.5 Accessing User Memory
 - It also gives you a solution: check whether the user-provided pointer belongs to user space (address below 3GB)

3.1.5 Accessing User Memory

As part of a system call, the kernel must often access memory through pointers provided by a user program. The kernel must be very careful about doing so, because the user can pass a null pointer, a pointer to unmapped virtual memory, or a pointer to kernel virtual address space (above `PHYS_BASE`). All of these types of invalid pointers must be rejected

The second method is to check only that a user pointer points below `PHYS_BASE`, then dereference it. An invalid user pointer will cause a “page fault” that you can handle by

Real-world Example

■ CVE-2020-0792 found in *Windows* kernel

- System call handler performs a series of check on user inputs
- `IsKernelSpace()` checks the range of a pointer
- But this check can be skipped if odd-number length is provided

```
1  SyscallHandler(struct Str* arg) {  
2      wchar_t *buf = arg->buf;  
3      ushort len = arg->len;  
4      if (len & 1) {  
5          LogError(...);  
6      }  
7      else if (IsKernelSpace(buf)) { // Why "else-if"?  
8          return;  
9      }  
10     ... // Access 'buf' here.
```

Quiz

■ Consider the following system call handler in kernel

- At first glance, this code seems to be safe
- But there is a potential vulnerability here: what is it?


```
syscallHandler(struct Str* arg) {  
    char kbuf[128];  
    if (IsKernelSpace(arg->buf) || arg->len >= 128) {  
        return;  
    }  
    memcpy(kbuf, arg->buf, arg->len);  
    ...  
}
```

Double Fetch (Race Condition)

■ The code fetches the buf (and len) field twice

- `arg->buf` or `arg->len` may change between the TOC and TOU
- Assume that an attacker creates two threads **t1** and **t2**
 - First, **t1** invokes system call and passes **IsKernelSpace()**

```
if (IsKernelSpace(arg->buf) || arg->len >= 128) {  
    return;  
}  
memcpy(kbuf, arg->buf, arg->len);
```




Double Fetch (Race Condition)

■ The code fetches the buf (and len) field twice

- `arg->buf` or `arg->len` may change between the TOC and TOU
- Assume that an attacker creates two threads **t1** and **t2**
 - First, **t1** invokes system call and passes **IsKernelSpace()**
 - Then the execution switches to thread **t2**, which updates **arg**
 - Finally, **t1** resumes and uses the updated field in **arg**

```
if (IsKernelSpace(arg->buf) || arg->len >= 128) {  
    return;  
}  
memcpy(kbuf, arg->buf, arg->len);
```



Preventing Double Fetch

- How can we avoid double fetch bug?
- Copy the user-provided inputs to the kernel space first (e.g., to local variables), and then perform validation
 - Now, the attacker cannot change the value of `buf` and `len` after passing the validation (TOCTOU not possible anymore)

```
SyscallHandler(struct Str* arg) {  
    char kbuf[128];  
    char *buf = arg->buf; // Copy arguments to local variable  
    int len = arg->len;  
    if (IsKernelSpace(buf) || len >= 128) { // Validation  
        return;  
    }  
    memcpy(kbuf, buf, len);  
    ...  
}
```


NULL Dereference in User Code

- Now, let's forget about the kernel for a while and think about this question:
 - In **user code**, can NULL dereference raise a security issue?
- If your application uses NULL pointer, it will crash
 - So this may result in a denial-of-service attack (if the application was running an important service)
- Can hackers do something more than denial-of-service?
 - In most cases, NULL dereference cannot lead to arbitrary code execution or memory disclosure
 - NULL is just an invalid memory address, so using such pointer will immediately raise a crash


NULL Dereference in Kernel

- NULL dereference **in kernel code** a totally different story
- Assume that the kernel code is doing certain operation
 - Ex) Writing a secret key value to a pointer
 - Ex) Fetching and calling a function pointer
 - What happens if these pointers (**p**, **fptr_table[i]**) are NULL?
 - Note that these pointers are **not** provided from user (via syscall)

```
int x, *p = NULL;
if (...) {
    p = &x;
}
*p = secret_key; // What if p is NULL?
...
fptr_table[i](100); // What if fptr_table[i] is NULL?
```

NULL Pointer \in User Space

- In kernel security, NULL pointer is not just an *invalid address*: it is an **address that belongs to the user space**
 - What if the user can allocate memory at address 0 (zero)?
 - You cannot do this with `malloc()`, but it is possible with `mmap()`
 - To be precise, it **was** possible (details in the next page)
 - Then the kernel will store the secret key to address 0 (user space)
 - Or the kernel may even execute the code in the user space



Preventing NULL Dereference

■ At software-level

- Prevents memory allocation at address 0
- Configuration parameter `mmap_min_addr`
 - User cannot allocate memory at address below this range

■ At hardware-level

- SMEP (Supervisor Mode **Execution** Prevention): prevents CPU running in the kernel mode from **executing** user-space code
- SMAP (Supervisor Mode **Access** Prevention): prevents CPU running in the kernel mode from **accessing** user-space data
 - Temporarily disabled when a system call handler has to fetch data from the user space memory

Lessons

- **The complex threat model of kernel security gives a rise to new interesting types of vulnerabilities**
 - Ex) NULL dereference is not a serious vulnerability in user application code, but it is a critical vulnerability in kernel code
- **Understanding certain kind of security issues requires a deep understanding on the internals of computer**
 - Ex) To understand the memory disclosure due to alignment, you should even know the behavior of a compiler

Appendix:

1 more subtle kernel vulnerability

Quiz

■ Consider the following system call handler in kernel

- At first glance, this code seems to be safe
- But there is a potential vulnerability here: what is it?

```
gettime_handler(struct Time *arg) {  
    struct Time t;  
    if (is_kernel_space(arg)) {  
        return;  
    }  
    // Set each field of the local struct  
    t.second = ...;  
    t.nano_sec = ...;  
    memcpy(arg, &t, sizeof(struct Time));  
}
```

```
struct Time {  
    int second;  
    long nano_sec;  
};
```


Review: Alignment

■ Assume a data type that requires K bytes of memory

- "Aligned" means that its memory address is multiple of K
- Ex) In order to be aligned, `int` type (4-byte) variable must be placed at an address that is multiple of 4

■ Motivation for alignment

- At hardware level, memory is accessed by chunk of bytes (word)
- Inefficient to load or store data that spans multiple words



Review: Alignment of Structure

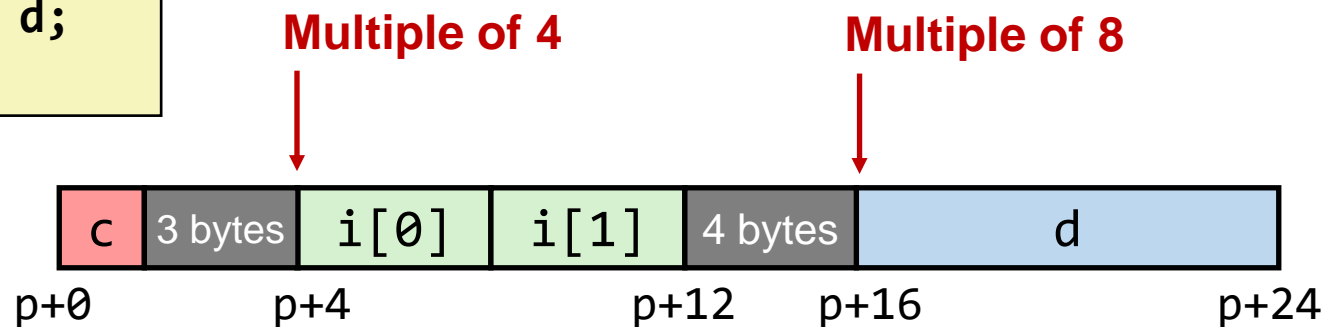
■ Alignment within the structure

- Each *field* must be placed at an *aligned offset*

■ Example: `struct S1`

- Compiler will insert *padding* (unused space) between the fields

```
struct S1 {  
    char c;  
    int i[2];  
    double d;  
};
```



Alignment and Memory Disclosure

■ Let's return to the example code in the previous quiz

- The padding bytes between the two fields are not initialized
- Uninitialized bytes in the stack will be copied to user space
- If unlucky, it will disclose sensitive data in kernel memory

```
gettime_handler(struct Time *arg) {  
    struct Time t;  
    if (is_kernel_space(arg)) {  
        return;  
    }  
    // Set each field of the local struct  
    t.second = ...;  
    t.nano_sec = ...;  
    memcpy(arg, &t, sizeof(struct Time));  
}
```

```
struct Time {  
    int second;  
    long nano_sec;  
};
```